

# DeepLearning\_HW1\_AmirPourmand

March 12, 2022

```
[1]: # Amir Pourmand - 99210259
```

```
[10]: # Amir Pourmand (99210259)

import torchvision.datasets as ds
from sklearn.utils import shuffle
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

data_train = np.array(ds.FashionMNIST(root="./data", train=True, download=True).
    ↳data)
target_train = np.array(ds.FashionMNIST(root="./data", train=True,
    ↳download=True).targets)
data_test = np.array(ds.FashionMNIST(root="./data", train=False, download=True).
    ↳data)
target_test = np.array(ds.FashionMNIST(root="./data", train=False,
    ↳download=True).targets)

#data_train, target_train = shuffle(data_train, target_train)

#### Transform the data! ####
data_train = data_train / 255
data_test = data_test / 255

target_train = pd.get_dummies(target_train).values
target_test = pd.get_dummies(target_test).values

data_train = data_train.reshape((-1, 28*28))
data_test = data_test.reshape((-1, 28*28))
```

```
[3]: from IPython.core.debugger import set_trace
```

#Part 1


```
[17]: def show_grid_twobytwo(data):
        my_pic = data[np.random.randint(0, 60000, 4)]
```

```

fig, axs = plt.subplots(2, 2)
axs[0, 0].imshow(my_pic[0].reshape(28,28))
axs[0, 1].imshow(my_pic[1].reshape(28,28))
axs[1, 0].imshow(my_pic[2].reshape(28,28))
axs[1, 1].imshow(my_pic[3].reshape(28,28))

show_grid_twobytwo(data_train)

```



#Part 2

```

[176]: class ReLU():
        def forward(self, x):
            self.old_x = np.copy(x)
            return np.clip(x,0,None)

        def backward(self, grad):
            return np.where(self.old_x>0,grad,0)

class Sigmoid():
    def forward(self, x):
        self.old_y = np.exp(x) / (1. + np.exp(x))
        return self.old_y

    def backward(self, grad):
        return self.old_y * (1. - self.old_y) * grad

```

```

class Softmax():
    def forward(self,x):
        self.old_y = np.exp(x) / np.exp(x).sum(axis=1)[:,None]
        return self.old_y

    def backward(self,grad):
        return self.old_y * (grad -(grad * self.old_y).sum(axis=1)[:,None])

class CrossEntropy():
    def forward(self,x,y):
        self.old_x = x.clip(min=1e-8,max=None)
        self.old_y = y
        return (np.where(y==1,-np.log(self.old_x), 0)).sum(axis=1)

    def backward(self):
        return np.where(self.old_y==1,-1/self.old_x, 0)

class Linear():
    def __init__(self,n_in,n_out,momentum=False):
        self.weights = np.random.randn(n_in,n_out) * np.sqrt(2/n_in)
        self.biases = np.zeros(n_out)

        self.momentum = momentum
        if momentum:
            self.grad_w =self.previous = np.zeros((n_in,n_out))

    def forward(self, x):
        self.old_x = x
        return np.dot(x,self.weights) + self.biases

    def backward(self,grad):
        self.grad_b = grad.mean(axis=0)

        if self.momentum:
            self.previous = self.grad_w

        self.grad_w = (np.matmul(self.old_x[:, :,None],grad[:,None,:])).
        ↪mean(axis=0)
        return np.dot(grad,self.weights.transpose())

```

```

[136]: class Model():
        def __init__(self, layers, cost):
            self.layers = layers
            self.cost = cost

        def forward(self,x):

```

```

        for layer in self.layers:
            x = layer.forward(x)
        return x

    def loss(self,x,y):
        self.output = self.forward(x)
        return self.cost.forward(self.output,y)

    def backward(self):
        grad = self.cost.backward()
        for i in range(len(self.layers)-1,-1,-1):
            grad = self.layers[i].backward(grad)

    def calculate_accuracy(self, real):
        predictedNumber = np.argmax(self.output,axis=1)
        realNumber = np.argmax(real,axis=1)
        return np.mean(predictedNumber == realNumber)

def train(model,lr,nb_epoch,data,label,batch_size=1000):
    loss_train = []
    loss_test = []

    accuracy_train = []
    accuracy_test = []
    for epoch in range(nb_epoch):
        lr = lr * 0.99
        if lr < 1e-3:
            lr = 1e-3

        num_inputs = 0
        running_loss = 0
        accuracy = 0

        number_of_batches = int(data.shape[0]/batch_size)
        for i in range(number_of_batches):
            input = data[i*batch_size:(i+1)*batch_size]
            output = label[i*batch_size:(i+1)*batch_size]
            num_inputs += input.shape[0]
            running_loss += model.loss(input,output).sum()
            model.backward()

            accuracy += model.calculate_accuracy(output)

        for layer in model.layers:
            if type(layer) == Linear:
                layer.weights -= lr * layer.grad_w

```

```

        layer.biases -= lr * layer.grad_b
    print(f'Epoch {epoch+1}/{nb_epoch}: loss train = {running_loss}')
    print(f'Train accuracy: {accuracy/number_of_batches}')

    loss_test_current = model.loss(data_test,target_test).sum()
    print(f'Epoch {epoch+1}/{nb_epoch}: loss test = {loss_test_current}')
    accuracy_current_test= model.calculate_accuracy(target_test)
    print(f'Test accuracy: {accuracy_current_test}')

    loss_train.append(running_loss)
    loss_test.append(loss_test_current)

    accuracy_train.append(accuracy/number_of_batches)
    accuracy_test.append(accuracy_current_test)
    return loss_train,loss_test,accuracy_train,accuracy_test

```

```

[139]: net = Model([Linear(784,128), Sigmoid(), Linear(128,64), Sigmoid(),
↳Linear(64,10), Softmax()], CrossEntropy())

```

```

[140]: epoch_count = 30
    loss_train,loss_test,accuracy_train,accuracy_test =
↳train(net,1e-1,epoch_count,data_train,target_train)

```

```

Epoch 1/30: loss train = 132827.31534895362
Train accuracy: 0.35323333333333334
Epoch 1/30: loss test = 20946.296056284016
Test accuracy: 0.4632
Epoch 2/30: loss train = 117172.60599323544
Train accuracy: 0.5205333333333333
Epoch 2/30: loss test = 17953.662677965036
Test accuracy: 0.5687
Epoch 3/30: loss train = 98626.18946359388
Train accuracy: 0.6223333333333333
Epoch 3/30: loss test = 15075.571274081087
Test accuracy: 0.66
Epoch 4/30: loss train = 83943.86310049995
Train accuracy: 0.66669999999999998
Epoch 4/30: loss test = 13095.937413259353
Test accuracy: 0.6757
Epoch 5/30: loss train = 73973.71860094678
Train accuracy: 0.68276666666666665
Epoch 5/30: loss test = 11739.020457773331
Test accuracy: 0.6895
Epoch 6/30: loss train = 66933.03985993215
Train accuracy: 0.69555000000000002
Epoch 6/30: loss test = 10751.809601783323
Test accuracy: 0.6981


```

Epoch 7/30: loss train = 61667.876132650716  
Train accuracy: 0.7071833333333334  
Epoch 7/30: loss test = 9995.114571604656  
Test accuracy: 0.7081  
Epoch 8/30: loss train = 57547.6836181342  
Train accuracy: 0.7169500000000002  
Epoch 8/30: loss test = 9392.592953102454  
Test accuracy: 0.7145  
Epoch 9/30: loss train = 54219.690978905935  
Train accuracy: 0.7259166666666667  
Epoch 9/30: loss test = 8900.556119154262  
Test accuracy: 0.7213  
Epoch 10/30: loss train = 51476.333175961896  
Train accuracy: 0.7324499999999997  
Epoch 10/30: loss test = 8492.417457977645  
Test accuracy: 0.7266  
Epoch 11/30: loss train = 49186.19135391677  
Train accuracy: 0.7386833333333331  
Epoch 11/30: loss test = 8150.55075674045  
Test accuracy: 0.7305  
Epoch 12/30: loss train = 47258.16462239753  
Train accuracy: 0.74345  
Epoch 12/30: loss test = 7862.141325124827  
Test accuracy: 0.7352  
Epoch 13/30: loss train = 45623.69243591608  
Train accuracy: 0.7477166666666668  
Epoch 13/30: loss test = 7617.196853815864  
Test accuracy: 0.7397  
Epoch 14/30: loss train = 44228.35205687464  
Train accuracy: 0.7510666666666668  
Epoch 14/30: loss test = 7407.643638566689  
Test accuracy: 0.7443  
Epoch 15/30: loss train = 43027.944823165264  
Train accuracy: 0.7541333333333332  
Epoch 15/30: loss test = 7226.9017491708255  
Test accuracy: 0.7486  
Epoch 16/30: loss train = 41986.448992140344  
Train accuracy: 0.7577333333333333  
Epoch 16/30: loss test = 7069.62804951542  
Test accuracy: 0.7513  
Epoch 17/30: loss train = 41074.62375804611  
Train accuracy: 0.7610500000000001  
Epoch 17/30: loss test = 6931.510152426372  
Test accuracy: 0.7538  
Epoch 18/30: loss train = 40268.860103027044  
Train accuracy: 0.7637333333333335  
Epoch 18/30: loss test = 6809.085700227674  
Test accuracy: 0.7569

Epoch 19/30: loss train = 39550.18896524748  
Train accuracy: 0.7662666666666667  
Epoch 19/30: loss test = 6699.585315690341  
Test accuracy: 0.7591  
Epoch 20/30: loss train = 38903.42894354668  
Train accuracy: 0.7691666666666666  
Epoch 20/30: loss test = 6600.79917734809  
Test accuracy: 0.761  
Epoch 21/30: loss train = 38316.4627883525  
Train accuracy: 0.7717000000000003  
Epoch 21/30: loss test = 6510.965634710272  
Test accuracy: 0.763  
Epoch 22/30: loss train = 37779.63061594186  
Train accuracy: 0.7742500000000001  
Epoch 22/30: loss test = 6428.67942762729  
Test accuracy: 0.7656  
Epoch 23/30: loss train = 37285.2270813171  
Train accuracy: 0.7766166666666667  
Epoch 23/30: loss test = 6352.816926654394  
Test accuracy: 0.7678  
Epoch 24/30: loss train = 36827.09009696092  
Train accuracy: 0.7792333333333338  
Epoch 24/30: loss test = 6282.476002181724  
Test accuracy: 0.7703  
Epoch 25/30: loss train = 36400.26954895911  
Train accuracy: 0.7815499999999999  
Epoch 25/30: loss test = 6216.928352809768  
Test accuracy: 0.7715  
Epoch 26/30: loss train = 36000.76493478377  
Train accuracy: 0.7839000000000002  
Epoch 26/30: loss test = 6155.582225125222  
Test accuracy: 0.7738  
Epoch 27/30: loss train = 35625.32065576926  
Train accuracy: 0.7856333333333333  
Epoch 27/30: loss test = 6097.953503577149  
Test accuracy: 0.7769  
Epoch 28/30: loss train = 35271.26760508346  
Train accuracy: 0.7879166666666667  
Epoch 28/30: loss test = 6043.643295294777  
Test accuracy: 0.7789  
Epoch 29/30: loss train = 34936.400542759235  
Train accuracy: 0.7897666666666665  
Epoch 29/30: loss test = 5992.320432960774  
Test accuracy: 0.7804  
Epoch 30/30: loss train = 34618.88255533996  
Train accuracy: 0.7919833333333334  
Epoch 30/30: loss test = 5943.707690938072  
Test accuracy: 0.7822


```
[141]: plt.plot(np.arange(epoch_count), loss_train)
```

```
[141]: [<matplotlib.lines.Line2D at 0x7f6f1290e550>]
```



```
[142]: plt.plot(np.arange(epoch_count), loss_test)
```


```
[142]: [<matplotlib.lines.Line2D at 0x7f6f128e6ed0>]
```






```
[143]: plt.plot(np.arange(epoch_count), accuracy_train)
```

```
[143]: [<matplotlib.lines.Line2D at 0x7f6f12856c90>]
```



```
[144]: plt.plot(np.arange(epoch_count), accuracy_test)
```

```
[144]: [<matplotlib.lines.Line2D at 0x7f6f127c3450>]
```



## 1 Part 3

```
[146]: net2 = Model([Linear(784,128), ReLU(), Linear(128,64), ReLU(), Linear(64,10),  
↪Softmax()], CrossEntropy())
```

```
[147]: epoch_count = 30  
        loss_train,loss_test,accuracy_train,accuracy_test =  
↪train(net,1e-1,epoch_count,data_train,target_train)
```

```
Epoch 1/30: loss train = 34273.999245328014  
Train accuracy: 0.7943333333333332  
Epoch 1/30: loss test = 5882.024706903974  
Test accuracy: 0.7845  
Epoch 2/30: loss train = 33891.26773555227  
Train accuracy: 0.7967333333333333  
Epoch 2/30: loss test = 5823.962497812203  
Test accuracy: 0.7865  
Epoch 3/30: loss train = 33530.60173532922  
Train accuracy: 0.7998333333333332  
Epoch 3/30: loss test = 5769.388843975354  
Test accuracy: 0.7884  
Epoch 4/30: loss train = 33190.318871548116  
Train accuracy: 0.8021333333333331  
Epoch 4/30: loss test = 5717.999626335039
```


Test accuracy: 0.7909  
Epoch 5/30: loss train = 32868.717160099586  
Train accuracy: 0.8039666666666668  
Epoch 5/30: loss test = 5669.53330350921  
Test accuracy: 0.793  
Epoch 6/30: loss train = 32564.331491516987  
Train accuracy: 0.8061499999999999  
Epoch 6/30: loss test = 5623.76092580437  
Test accuracy: 0.7947  
Epoch 7/30: loss train = 32275.88105378048  
Train accuracy: 0.8082166666666668  
Epoch 7/30: loss test = 5580.479400607843  
Test accuracy: 0.7966  
Epoch 8/30: loss train = 32002.23148273933  
Train accuracy: 0.8099833333333335  
Epoch 8/30: loss test = 5539.50667765365  
Test accuracy: 0.7981  
Epoch 9/30: loss train = 31742.367252942953  
Train accuracy: 0.8119666666666666  
Epoch 9/30: loss test = 5500.6782600567085  
Test accuracy: 0.7987  
Epoch 10/30: loss train = 31495.371060447782  
Train accuracy: 0.8138166666666667  
Epoch 10/30: loss test = 5463.844607662446  
Test accuracy: 0.8002  
Epoch 11/30: loss train = 31260.40798246758  
Train accuracy: 0.8151333333333333  
Epoch 11/30: loss test = 5428.869133093598  
Test accuracy: 0.8016  
Epoch 12/30: loss train = 31036.71298637588  
Train accuracy: 0.8166833333333335  
Epoch 12/30: loss test = 5395.626598800164  
Test accuracy: 0.8035  
Epoch 13/30: loss train = 30823.580907713233  
Train accuracy: 0.8181833333333334  
Epoch 13/30: loss test = 5364.001799454714  
Test accuracy: 0.805  
Epoch 14/30: loss train = 30620.35836044925  
Train accuracy: 0.8192999999999998  
Epoch 14/30: loss test = 5333.888461616467  
Test accuracy: 0.8055  
Epoch 15/30: loss train = 30426.437234558463  
Train accuracy: 0.8205166666666667  
Epoch 15/30: loss test = 5305.1883190221915  
Test accuracy: 0.807  
Epoch 16/30: loss train = 30241.2495282917  
Train accuracy: 0.8218333333333335  
Epoch 16/30: loss test = 5277.8103347078195

Test accuracy: 0.8079  
Epoch 17/30: loss train = 30064.263299206854  
Train accuracy: 0.8232166666666667  
Epoch 17/30: loss test = 5251.670046556537  
Test accuracy: 0.8094  
Epoch 18/30: loss train = 29894.979530694858  
Train accuracy: 0.8239833333333337  
Epoch 18/30: loss test = 5226.689015030968  
Test accuracy: 0.8109  
Epoch 19/30: loss train = 29732.92971884649  
Train accuracy: 0.8248666666666669  
Epoch 19/30: loss test = 5202.794353242951  
Test accuracy: 0.8112  
Epoch 20/30: loss train = 29577.67399753661  
Train accuracy: 0.8257000000000002  
Epoch 20/30: loss test = 5179.918321275403  
Test accuracy: 0.8126  
Epoch 21/30: loss train = 29428.79964011735  
Train accuracy: 0.8265333333333333  
Epoch 21/30: loss test = 5157.997969074271  
Test accuracy: 0.8135  
Epoch 22/30: loss train = 29285.91980267913  
Train accuracy: 0.8275833333333334  
Epoch 22/30: loss test = 5136.974815103587  
Test accuracy: 0.8142  
Epoch 23/30: loss train = 29148.672403354012  
Train accuracy: 0.8286333333333337  
Epoch 23/30: loss test = 5116.794550974292  
Test accuracy: 0.8152  
Epoch 24/30: loss train = 29016.719061362473  
Train accuracy: 0.8293833333333336  
Epoch 24/30: loss test = 5097.406765101522  
Test accuracy: 0.8161  
Epoch 25/30: loss train = 28889.74404589522  
Train accuracy: 0.8299833333333335  
Epoch 25/30: loss test = 5078.764680891482  
Test accuracy: 0.8163  
Epoch 26/30: loss train = 28767.453206914564  
Train accuracy: 0.8308833333333336  
Epoch 26/30: loss test = 5060.824906893904  
Test accuracy: 0.8172  
Epoch 27/30: loss train = 28649.572876927738  
Train accuracy: 0.8312833333333335  
Epoch 27/30: loss test = 5043.547197759708  
Test accuracy: 0.8169  
Epoch 28/30: loss train = 28535.848744823237  
Train accuracy: 0.8321999999999999  
Epoch 28/30: loss test = 5026.894225762419

```
Test accuracy: 0.817
Epoch 29/30: loss train = 28426.044710542068
Train accuracy: 0.83300000000000002
Epoch 29/30: loss test = 5010.831363159079
Test accuracy: 0.8176
Epoch 30/30: loss train = 28319.941733479172
Train accuracy: 0.8335833333333335
Epoch 30/30: loss test = 4995.326475878811
Test accuracy: 0.818
```


```
[148]: plt.plot(np.arange(epoch_count), loss_train)
```

```
[148]: [<matplotlib.lines.Line2D at 0x7f6f12764c10>]
```




```
[149]: plt.plot(np.arange(epoch_count), loss_test)
```

```
[149]: [<matplotlib.lines.Line2D at 0x7f6f126d4d10>]
```




```
[150]: plt.plot(np.arange(epoch_count), accuracy_train)
```

```
[150]: [<matplotlib.lines.Line2D at 0x7f6f1263d690>]
```



```
[151]: plt.plot(np.arange(epoch_count), accuracy_test)
```

```
[151]: [<matplotlib.lines.Line2D at 0x7f6f125b1c50>]
```



#Part 4

```
[179]: class Model():
    def __init__(self, layers, cost):
        self.layers = layers
        self.cost = cost

    def forward(self, x):
        for layer in self.layers:
            x = layer.forward(x)
        return x

    def loss(self, x, y):
        self.output = self.forward(x)
        return self.cost.forward(self.output, y)

    def backward(self):
        grad = self.cost.backward()
        for i in range(len(self.layers)-1, -1, -1):
            grad = self.layers[i].backward(grad)

    def step(self, lr):
```

```

        for layer in self.layers:
            if type(layer) == Linear:
                layer.previous = lr* layer.grad_w +momentum_rate*layer.
→previous

                layer.weights -= layer.previous
                layer.biases -= lr * layer.grad_b

def calculate_accuracy(self, real):
    predictedNumber = np.argmax(self.output,axis=1)
    realNumber = np.argmax(real,axis=1)
    return np.mean(predictedNumber == realNumber)

def train_momentum(model,lr,nb_epoch,data,label,batch_size=1000):
    loss_train = []
    loss_test = []

    accuracy_train = []
    accuracy_test = []
    for epoch in range(nb_epoch):
        lr = lr * 0.99
        if lr < 1e-3:
            lr = 1e-3

        num_inputs = 0
        running_loss = 0
        accuracy = 0

        number_of_batches = int(data.shape[0]/batch_size)
        for i in range(number_of_batches):
            input = data[i*batch_size:(i+1)*batch_size]
            output = label[i*batch_size:(i+1)*batch_size]
            num_inputs += input.shape[0]
            running_loss += model.loss(input,output).sum()
            model.backward()

            accuracy += model.calculate_accuracy(output)
            model.step(lr)

        print(f'Epoch {epoch+1}/{nb_epoch}: loss train = {running_loss}')
        print(f'Train accuracy: {accuracy/number_of_batches}')

        loss_test_current = model.loss(data_test,target_test).sum()
        print(f'Epoch {epoch+1}/{nb_epoch}: loss test = {loss_test_current}')
        accuracy_current_test= model.calculate_accuracy(target_test)
        print(f'Test accuracy: {accuracy_current_test}')

```



```

        loss_train.append(running_loss)
        loss_test.append(loss_test_current)

        accuracy_train.append(accuracy/number_of_batches)
        accuracy_test.append(accuracy_current_test)
    return loss_train,loss_test,accuracy_train,accuracy_test

```

```

[180]: net = Model([Linear(784,128,True), Sigmoid(), Linear(128,64,True), Sigmoid(),
    ↳Linear(64,10,True), Softmax()], CrossEntropy())
epoch_count = 30
loss_train,loss_test,accuracy_train,accuracy_test =
    ↳train_momentum(net,1e-1,epoch_count,data_train,target_train)

```

```

Epoch 1/30: loss train = 103761.94796602237
Train accuracy: 0.37478333333333334
Epoch 1/30: loss test = 11689.498285159467
Test accuracy: 0.538
Epoch 2/30: loss train = 58272.40424204931
Train accuracy: 0.6165
Epoch 2/30: loss test = 9527.279536904098
Test accuracy: 0.6295
Epoch 3/30: loss train = 46150.98579169021
Train accuracy: 0.69663333333333334
Epoch 3/30: loss test = 7587.40233664309
Test accuracy: 0.7029
Epoch 4/30: loss train = 41056.18575142075
Train accuracy: 0.73146666666666665
Epoch 4/30: loss test = 7156.556337217709
Test accuracy: 0.729
Epoch 5/30: loss train = 38149.82667771503
Train accuracy: 0.75108333333333333
Epoch 5/30: loss test = 6572.531684206939
Test accuracy: 0.7462
Epoch 6/30: loss train = 35604.57046914616
Train accuracy: 0.77133333333333338
Epoch 6/30: loss test = 5814.691801247
Test accuracy: 0.7793
Epoch 7/30: loss train = 34009.61060029159
Train accuracy: 0.78336666666666667
Epoch 7/30: loss test = 5458.487971309901
Test accuracy: 0.7955
Epoch 8/30: loss train = 32813.69166456551
Train accuracy: 0.79463333333333334
Epoch 8/30: loss test = 6009.993532053346
Test accuracy: 0.7773
Epoch 9/30: loss train = 32275.19084573284
Train accuracy: 0.79809999999999998


```

Epoch 9/30: loss train = 5226.966779154435  
Test accuracy: 0.8034  
Epoch 10/30: loss train = 31000.030595849257  
Train accuracy: 0.8076666666666665  
Epoch 10/30: loss test = 5726.385985353482  
Test accuracy: 0.7803  
Epoch 11/30: loss train = 29855.372698577157  
Train accuracy: 0.8162333333333334  
Epoch 11/30: loss test = 5846.988979506997  
Test accuracy: 0.7754  
Epoch 12/30: loss train = 29300.90003255  
Train accuracy: 0.8196166666666667  
Epoch 12/30: loss test = 5221.322454483814  
Test accuracy: 0.8049  
Epoch 13/30: loss train = 29285.494567831876  
Train accuracy: 0.8199666666666665  
Epoch 13/30: loss test = 4990.5608097932545  
Test accuracy: 0.8124  
Epoch 14/30: loss train = 28366.30472880257  
Train accuracy: 0.8259333333333335  
Epoch 14/30: loss test = 4888.639831028579  
Test accuracy: 0.8213  
Epoch 15/30: loss train = 27717.817467901965  
Train accuracy: 0.8301666666666668  
Epoch 15/30: loss test = 4860.8972437192215  
Test accuracy: 0.8182  
Epoch 16/30: loss train = 27116.76668743817  
Train accuracy: 0.8342500000000003  
Epoch 16/30: loss test = 4612.4464375894395  
Test accuracy: 0.8329  
Epoch 17/30: loss train = 26390.402952827724  
Train accuracy: 0.8392666666666666  
Epoch 17/30: loss test = 4596.484309709349  
Test accuracy: 0.8337  
Epoch 18/30: loss train = 26068.86815108737  
Train accuracy: 0.8413333333333334  
Epoch 18/30: loss test = 4539.911092383581  
Test accuracy: 0.8351  
Epoch 19/30: loss train = 25661.87665136205  
Train accuracy: 0.8440999999999999  
Epoch 19/30: loss test = 4504.444635118725  
Test accuracy: 0.8355  
Epoch 20/30: loss train = 25340.22516857293  
Train accuracy: 0.8462166666666664  
Epoch 20/30: loss test = 4430.968041304  
Test accuracy: 0.8404  
Epoch 21/30: loss train = 25984.311502511213  
Train accuracy: 0.8413166666666665

```
Epoch 21/30: loss test = 4519.952104775767
Test accuracy: 0.832
Epoch 22/30: loss train = 25441.60161612542
Train accuracy: 0.8435833333333334
Epoch 22/30: loss test = 4340.280456569409
Test accuracy: 0.8414
Epoch 23/30: loss train = 24852.244826296108
Train accuracy: 0.8480333333333332
Epoch 23/30: loss test = 4491.007500768186
Test accuracy: 0.8311
Epoch 24/30: loss train = 24582.163285242055
Train accuracy: 0.8493333333333334
Epoch 24/30: loss test = 4348.987067474437
Test accuracy: 0.8406
Epoch 25/30: loss train = 24477.176135132024
Train accuracy: 0.85015
Epoch 25/30: loss test = 4510.567816994593
Test accuracy: 0.8295
Epoch 26/30: loss train = 24054.759123205233
Train accuracy: 0.8529833333333331
Epoch 26/30: loss test = 4303.107095474756
Test accuracy: 0.8426
Epoch 27/30: loss train = 23906.170405875644
Train accuracy: 0.8536166666666664
Epoch 27/30: loss test = 4419.326099768561
Test accuracy: 0.8321
Epoch 28/30: loss train = 23537.280939203625
Train accuracy: 0.8561166666666665
Epoch 28/30: loss test = 4261.2626447180155
Test accuracy: 0.8446
Epoch 29/30: loss train = 23469.761025458978
Train accuracy: 0.8566166666666667
Epoch 29/30: loss test = 4350.221229163766
Test accuracy: 0.8349
Epoch 30/30: loss train = 23102.820436305625
Train accuracy: 0.8586333333333334
Epoch 30/30: loss test = 4222.560625831047
Test accuracy: 0.8442
```


```
[181]: plt.plot(np.arange(epoch_count), loss_train)
```

```
[181]: [<matplotlib.lines.Line2D at 0x7f6f1b4ec690>]
```




```
[182]: plt.plot(np.arange(epoch_count), loss_test)
```

```
[182]: [<matplotlib.lines.Line2D at 0x7f6f1b425090>]
```




```
[183]: plt.plot(np.arange(epoch_count), accuracy_train)
```

```
[183]: [<matplotlib.lines.Line2D at 0x7f6f1b3b9810>]
```



```
[184]: plt.plot(np.arange(epoch_count), accuracy_test)
```

```
[184]: [<matplotlib.lines.Line2D at 0x7f6f121fb790>]
```



#Part 5

```
[158]: net = Model([Linear(784,128), Sigmoid(), Linear(128,64), Sigmoid(),  
↳Linear(64,64), Sigmoid(), Linear(64,10), Softmax()], CrossEntropy())  
epoch_count = 30  
loss_train,loss_test,accuracy_train,accuracy_test =  
↳train(net,1e-1,epoch_count,data_train,target_train)
```

```
Epoch 1/30: loss train = 138585.16783004766  
Train accuracy: 0.17660000000000003  
Epoch 1/30: loss test = 22762.64773987977  
Test accuracy: 0.2118  
Epoch 2/30: loss train = 135560.2365353218  
Train accuracy: 0.29771666666666667  
Epoch 2/30: loss test = 22382.228178991583  
Test accuracy: 0.3337  
Epoch 3/30: loss train = 132365.951542139  
Train accuracy: 0.34273333333333334  
Epoch 3/30: loss test = 21645.577227901616  
Test accuracy: 0.3456  
Epoch 4/30: loss train = 126037.53291557218  
Train accuracy: 0.36085000000000001  
Epoch 4/30: loss test = 20229.266243685433  
Test accuracy: 0.3469  
Epoch 5/30: loss train = 115819.46039339514  
Train accuracy: 0.38485  
Epoch 5/30: loss test = 18383.926078942066  
Test accuracy: 0.3973  
Epoch 6/30: loss train = 105550.65833693123  
Train accuracy: 0.42003333333333332  
Epoch 6/30: loss test = 16907.74045055039  
Test accuracy: 0.4585  
Epoch 7/30: loss train = 97911.1664760458  
Train accuracy: 0.460833333333333343  
Epoch 7/30: loss test = 15818.818293374112  
Test accuracy: 0.4981  
Epoch 8/30: loss train = 91937.08297539396  
Train accuracy: 0.50315  
Epoch 8/30: loss test = 14915.925281496035  
Test accuracy: 0.5319  
Epoch 9/30: loss train = 86844.825416397  
Train accuracy: 0.53845  
Epoch 9/30: loss test = 14134.816090919361  
Test accuracy: 0.5542  
Epoch 10/30: loss train = 82420.77233244957
```


Train accuracy: 0.5630499999999999  
Epoch 10/30: loss test = 13453.345091126657  
Test accuracy: 0.5694  
Epoch 11/30: loss train = 78543.79782535035  
Train accuracy: 0.5785666666666667  
Epoch 11/30: loss test = 12851.956481709045  
Test accuracy: 0.5802  
Epoch 12/30: loss train = 75113.73502583563  
Train accuracy: 0.5904333333333333  
Epoch 12/30: loss test = 12317.988602593487  
Test accuracy: 0.5897  
Epoch 13/30: loss train = 72077.86746367218  
Train accuracy: 0.5990333333333333  
Epoch 13/30: loss test = 11846.80038974688  
Test accuracy: 0.5971  
Epoch 14/30: loss train = 69414.31090835358  
Train accuracy: 0.6063166666666664  
Epoch 14/30: loss test = 11435.466236576582  
Test accuracy: 0.6046  
Epoch 15/30: loss train = 67093.61227173633  
Train accuracy: 0.6133500000000001  
Epoch 15/30: loss test = 11077.381948455151  
Test accuracy: 0.6127  
Epoch 16/30: loss train = 65063.677383283575  
Train accuracy: 0.6206833333333335  
Epoch 16/30: loss test = 10762.606858764037  
Test accuracy: 0.6184  
Epoch 17/30: loss train = 63262.05335185032  
Train accuracy: 0.6285333333333333  
Epoch 17/30: loss test = 10480.876838989785  
Test accuracy: 0.6267  
Epoch 18/30: loss train = 61631.878124954506  
Train accuracy: 0.6358833333333332  
Epoch 18/30: loss test = 10223.642812593505  
Test accuracy: 0.6336  
Epoch 19/30: loss train = 60129.12563321756  
Train accuracy: 0.6437333333333336  
Epoch 19/30: loss test = 9984.603493831346  
Test accuracy: 0.6402  
Epoch 20/30: loss train = 58722.97637278881  
Train accuracy: 0.6518833333333337  
Epoch 20/30: loss test = 9759.49189162154  
Test accuracy: 0.6515  
Epoch 21/30: loss train = 57393.5609399073  
Train accuracy: 0.6589999999999999  
Epoch 21/30: loss test = 9545.661567798636  
Test accuracy: 0.6595  
Epoch 22/30: loss train = 56129.1796259991

```
Train accuracy: 0.6665
Epoch 22/30: loss test = 9341.65331099526
Test accuracy: 0.668
Epoch 23/30: loss train = 54923.61518621588
Train accuracy: 0.6737666666666666
Epoch 23/30: loss test = 9146.779617459799
Test accuracy: 0.6754
Epoch 24/30: loss train = 53773.80243367801
Train accuracy: 0.6803666666666666
Epoch 24/30: loss test = 8960.76117081728
Test accuracy: 0.6849
Epoch 25/30: loss train = 52678.09744758627
Train accuracy: 0.6871666666666669
Epoch 25/30: loss test = 8783.46252907259
Test accuracy: 0.6954
Epoch 26/30: loss train = 51635.22966015136
Train accuracy: 0.6943833333333334
Epoch 26/30: loss test = 8614.743588202917
Test accuracy: 0.7007
Epoch 27/30: loss train = 50643.786053397365
Train accuracy: 0.7004666666666667
Epoch 27/30: loss test = 8454.39890020202
Test accuracy: 0.7074
Epoch 28/30: loss train = 49702.00319966678
Train accuracy: 0.7054499999999998
Epoch 28/30: loss test = 8302.145242952021
Test accuracy: 0.7131
Epoch 29/30: loss train = 48807.729850920936
Train accuracy: 0.7116499999999999
Epoch 29/30: loss test = 8157.633206340808
Test accuracy: 0.7186
Epoch 30/30: loss train = 47958.50888564728
Train accuracy: 0.7166999999999999
Epoch 30/30: loss test = 8020.472592466269
Test accuracy: 0.7217
```

```
[159]: plt.plot(np.arange(epoch_count), loss_train)
```


```
[159]: [<matplotlib.lines.Line2D at 0x7f6f12395610>]
```






```
[160]: plt.plot(np.arange(epoch_count), loss_test)
```

```
[160]: [<matplotlib.lines.Line2D at 0x7f6f122f9810>]
```




```
[161]: plt.plot(np.arange(epoch_count), accuracy_train)
```

```
[161]: [<matplotlib.lines.Line2D at 0x7f6f1282af50>]
```



```
[162]: plt.plot(np.arange(epoch_count), accuracy_test)
```

```
[162]: [<matplotlib.lines.Line2D at 0x7f6f125d3110>]
```



#Part 6

```
[218]: import numpy as np
import torch
import torchvision
import matplotlib.pyplot as plt
from time import time
from torchvision import datasets, transforms
from torch import nn, optim
```

```
[219]: transform = transforms.Compose([transforms.ToTensor(),
                                     transforms.Normalize((0.5,), (0.5,)),
                                     ])

trainset = datasets.FashionMNIST(root="./data", download=True, train=True,
    ↪transform=transform)
valset = datasets.FashionMNIST(root="./data", download=True, train=False,
    ↪transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=64, shuffle=True)
valloader = torch.utils.data.DataLoader(valset, batch_size=64, shuffle=True)

input_size = 784
hidden_sizes = [128, 64]
output_size = 10

model = nn.Sequential(nn.Linear(input_size, hidden_sizes[0]),
                      nn.Sigmoid(),
                      nn.Linear(hidden_sizes[0], hidden_sizes[1]),
                      nn.Sigmoid(),
                      nn.Linear(hidden_sizes[1], output_size),
                      nn.Softmax(dim=1))

print(model)
```

```
Sequential(
  (0): Linear(in_features=784, out_features=128, bias=True)
  (1): Sigmoid()
  (2): Linear(in_features=128, out_features=64, bias=True)
  (3): Sigmoid()
  (4): Linear(in_features=64, out_features=10, bias=True)
  (5): Softmax(dim=1)
)
```

```
[220]: criterion = nn.CrossEntropyLoss()
images, labels = next(iter(trainloader))
images = images.view(images.shape[0], -1)
```

```
logps = model(images)
loss = criterion(logps, labels)
```

```
[221]: optimizer = optim.SGD(model.parameters(), lr=0.03)
time0 = time()

loss_train = []
accuracy_train = []
accuracy_test = []

epochs = 15
for e in range(epochs):
    running_loss = 0
    for images, labels in trainloader:
        images = images.view(images.shape[0], -1)
        optimizer.zero_grad()
        output = model(images)
        loss = criterion(output, labels)

        loss.backward()

        optimizer.step()

        running_loss += loss.item()
    else:
        print("Epoch {} - Training loss: {}".format(e, running_loss/
→len(trainloader)))
        loss_train.append(running_loss)
        accuracy_train.append(calculate_accuracy(trainloader, 'train'))
        accuracy_test.append(calculate_accuracy(valloader, 'test'))

print("\nTraining Time (in minutes) =", (time()-time0)/60)
```

Epoch 0 - Training loss: 2.3021124784372002

Number Of Images Tested = 60000

Model Accuracy for train= 0.12763333333333332

Number Of Images Tested = 10000

Model Accuracy for test= 0.1302

Epoch 1 - Training loss: 2.300212318709156

Number Of Images Tested = 60000

Model Accuracy for train= 0.10668333333333334

Number Of Images Tested = 10000

Model Accuracy for test= 0.1069  
Epoch 2 - Training loss: 2.297341703860236  
Number Of Images Tested = 60000

Model Accuracy for train= 0.11051666666666667  
Number Of Images Tested = 10000

Model Accuracy for test= 0.1097  
Epoch 3 - Training loss: 2.2911456797931242  
Number Of Images Tested = 60000

Model Accuracy for train= 0.13371666666666668  
Number Of Images Tested = 10000

Model Accuracy for test= 0.1314  
Epoch 4 - Training loss: 2.2720074648541937  
Number Of Images Tested = 60000

Model Accuracy for train= 0.22395  
Number Of Images Tested = 10000

Model Accuracy for test= 0.222  
Epoch 5 - Training loss: 2.231057080124487  
Number Of Images Tested = 60000

Model Accuracy for train= 0.27018333333333333  
Number Of Images Tested = 10000

Model Accuracy for test= 0.2688  
Epoch 6 - Training loss: 2.186424291210134  
Number Of Images Tested = 60000

Model Accuracy for train= 0.28396666666666665  
Number Of Images Tested = 10000

Model Accuracy for test= 0.2843  
Epoch 7 - Training loss: 2.1435922488474897  
Number Of Images Tested = 60000

Model Accuracy for train= 0.34896666666666665  
Number Of Images Tested = 10000

Model Accuracy for test= 0.3498  
Epoch 8 - Training loss: 2.0989299757140025  
Number Of Images Tested = 60000

Model Accuracy for train= 0.44173333333333333  
Number Of Images Tested = 10000

Model Accuracy for test= 0.4409  
Epoch 9 - Training loss: 2.054632861095705  
Number Of Images Tested = 60000

Model Accuracy for train= 0.4580166666666667  
Number Of Images Tested = 10000

Model Accuracy for test= 0.4559  
Epoch 10 - Training loss: 2.016046718874974  
Number Of Images Tested = 60000

Model Accuracy for train= 0.5089166666666667  
Number Of Images Tested = 10000

Model Accuracy for test= 0.5086  
Epoch 11 - Training loss: 1.9852302403592352  
Number Of Images Tested = 60000

Model Accuracy for train= 0.5230333333333334  
Number Of Images Tested = 10000

Model Accuracy for test= 0.5248  
Epoch 12 - Training loss: 1.9615087304542314  
Number Of Images Tested = 60000

Model Accuracy for train= 0.5348833333333334  
Number Of Images Tested = 10000

Model Accuracy for test= 0.5356  
Epoch 13 - Training loss: 1.932801732630618  
Number Of Images Tested = 60000

Model Accuracy for train= 0.6055666666666667  
Number Of Images Tested = 10000

Model Accuracy for test= 0.605  
Epoch 14 - Training loss: 1.9018612763266574  
Number Of Images Tested = 60000

Model Accuracy for train= 0.6156333333333334  
Number Of Images Tested = 10000

Model Accuracy for test= 0.6143

Training Time (in minutes) = 6.454254253705343

```
[222]: def calculate_accuracy(loader,label):
        correct_count, all_count = 0, 0
        for images,labels in loader:
            for i in range(len(labels)):
                img = images[i].view(1, 784)
                with torch.no_grad():
                    logps = model(img)

                ps = torch.exp(logps)
                probab = list(ps.numpy()[0])
                pred_label = probab.index(max(probab))
                true_label = labels.numpy()[i]
                if(true_label == pred_label):
                    correct_count += 1
                all_count += 1

        print("Number Of Images Tested =", all_count)
        print("\nModel Accuracy for " + label + "=", (correct_count/all_count))
        return correct_count/all_count
```

[222]:

```
[223]: calculate_accuracy(valloader,'test')
```


Number Of Images Tested = 10000

Model Accuracy for test= 0.6143

[223]: 0.6143


```
[224]: plt.plot(np.arange(15), loss_train)
```

[224]: [ <matplotlib.lines.Line2D at 0x7f6f1130f650>]



```
[225]: plt.plot(np.arange(15), accuracy_train)
```


```
[225]: [<matplotlib.lines.Line2D at 0x7f6f1126a850>]
```





```
[226]: plt.plot(np.arange(15), accuracy_test)
```

```
[226]: [<matplotlib.lines.Line2D at 0x7f6f1115dd10>]
```



```
[226]:
```