

# DL\_HW2\_Exercise4

March 12, 2022

```
[1]: #Stu no: 99210259
     #stu name: Amir Pourmand
```

```
[16]: device = 'cuda' if torch.cuda.is_available() else 'cpu'
      print(f'Using {device} device')
```

Using cuda device

```
[25]: cuda = torch.device('cuda')
```

```
[2]: !wget https://github.com/mralisoltani/CNN_Tumor/raw/main/Tumor.zip
```

```
--2021-12-04 19:48:40--
https://github.com/mralisoltani/CNN_Tumor/raw/main/Tumor.zip
Resolving github.com (github.com)... 13.114.40.48
Connecting to github.com (github.com)|13.114.40.48|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location:
https://raw.githubusercontent.com/mralisoltani/CNN_Tumor/main/Tumor.zip
[following]
--2021-12-04 19:48:40--
https://raw.githubusercontent.com/mralisoltani/CNN_Tumor/main/Tumor.zip
Resolving raw.githubusercontent.com (raw.githubusercontent.com)...
185.199.108.133, 185.199.109.133, 185.199.111.133, ...
Connecting to raw.githubusercontent.com
(raw.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 14338568 (14M) [application/zip]
Saving to: 'Tumor.zip'

Tumor.zip          100%[=====>]  13.67M  --.-KB/s    in 0.1s

2021-12-04 19:48:42 (102 MB/s) - 'Tumor.zip' saved [14338568/14338568]
```

```
[ ]: !unzip /content/Tumor.zip
```

```
[ ]:
```

```
[74]: # -*- coding: utf-8 -*-
      """

      @author: Ali Soltani
      """

import torch
from torchvision import transforms
import os
from PIL import Image
import numpy as np
import pandas as pds
from torch.utils.data import DataLoader, TensorDataset, random_split

##### Loading Data
n = 3762
image=[]
cw = os.getcwd().replace(os.sep, '/')
trans = transforms.Compose([transforms.ToTensor()])
for i in range(n):
#     image.append(np.asarray(Image.open(cw + "/Brain_Tumor/Image" + str(i+1) + ".
    ↪ ".jpg"))))
    image.append(np.array(Image.open(cw + "/Brain_Tumor/Image" + str(i+1) + ".
    ↪ jpg").resize((48,48))))

temp = pds.read_csv(cw + "/Brain_Tumor.csv", index_col=None, header=None).
    ↪ to_numpy()
temp = temp[1:,1]
targets = np.zeros((n,1), dtype=int)
targets = []
for i in range(n):
    targets.append(int(temp[i]))

data = np.array(image)
data = data/255
data = torch.from_numpy(data).permute((0,3,2,1))
data = data.float().to(cuda)
targets = torch.tensor(targets).to(cuda)
dataset = TensorDataset(data, targets)
batch_size = 4
val_size = int(np.ceil(len(dataset)*0.2))
train_size = len(dataset) - val_size

train_data, test_data = random_split(dataset, [train_size, val_size])
```


```
train_loader = DataLoader(train_data, batch_size = batch_size, shuffle=True)
test_loader = DataLoader(test_data, batch_size = batch_size, shuffle=True)
```

```
[29]: import matplotlib.pyplot as plt
      from torchvision import transforms
```


```
[30]: print("count of data:", len(train_data))
      train_features, train_labels = next(iter(train_loader))
      print(f"Feature batch shape: {train_features.size()}")
      print(f"Labels batch shape: {train_labels.size()}")

      for i in range(batch_size):
          img = train_features[i].squeeze()
          label = train_labels[i]
          im = transforms.ToPILImage()(img).convert("RGB")
          plt.imshow(im)
          plt.show()
          print(f"Label: {label}")
```


```
count of data: 3009
Feature batch shape: torch.Size([4, 3, 48, 48])
Labels batch shape: torch.Size([4])
```




Label: 1



Label: 0



Label: 1



Label: 0

```
[7]: import os
import torch
from torch import nn
from torch.utils.data import DataLoader
from torchvision import datasets, transforms

class NeuralNetwork(nn.Module):
    def __init__(self):
        super(NeuralNetwork, self).__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(3*48*48, 512),
            nn.ReLU(),
            nn.Linear(512, 24),
            nn.ReLU(),
            nn.Linear(24, 2),
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits
```

```
[8]:
```

Using cuda device

```
[56]: model = NeuralNetwork().to(device)
      print(model)
```

```
NeuralNetwork(
  (flatten): Flatten(start_dim=1, end_dim=-1)
  (linear_relu_stack): Sequential(
    (0): Linear(in_features=6912, out_features=512, bias=True)
    (1): ReLU()
    (2): Linear(in_features=512, out_features=24, bias=True)
    (3): ReLU()
    (4): Linear(in_features=24, out_features=2, bias=True)
  )
)
```

```
[57]: learning_rate = 1e-3
      batch_size = 64
      epochs = 100

      loss_fn = nn.CrossEntropyLoss()

      optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
```

```
[58]: def train_loop(dataloader, model, loss_fn, optimizer):
      size = len(dataloader.dataset)
      for batch, (X, y) in enumerate(dataloader):
          pred = model(X)
          loss = loss_fn(pred, y)

          optimizer.zero_grad()
          loss.backward()
          optimizer.step()

          if batch % 100 == 0:
              loss, current = loss.item(), batch * len(X)
              print(f"loss: {loss:>7f} [{current:>5d}/{size:>5d}]")

      def test_loop(dataloader, model, loss_fn):
          size = len(dataloader.dataset)
          num_batches = len(dataloader)
          test_loss, correct = 0, 0

          with torch.no_grad():
              for X, y in dataloader:
                  pred = model(X)
```

```

        test_loss += loss_fn(pred, y).item()
        correct += (pred.argmax(1) == y).type(torch.float).sum().item()

    test_loss /= num_batches
    correct /= size
    print(f"Test Error: \n Accuracy: {(100*correct):>0.1f}%, Avg loss:␣
↪{test_loss:>8f} \n")
    return test_loss

```

```

[59]: loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

losses=[]
for t in range(epochs):
    print(f"Epoch {t+1}\n-----")
    train_loop(train_loader, model, loss_fn, optimizer)
    loss = test_loop(test_loader, model, loss_fn)
    losses.append(loss)
print("Done!")

```

Epoch 1

```

-----
loss: 0.612874 [ 0/ 3009]
loss: 0.712673 [ 400/ 3009]
loss: 0.744615 [ 800/ 3009]
loss: 0.690195 [1200/ 3009]
loss: 0.710857 [1600/ 3009]
loss: 0.616541 [2000/ 3009]
loss: 0.650407 [2400/ 3009]
loss: 0.670259 [2800/ 3009]
Test Error:
Accuracy: 78.8%, Avg loss: 0.609731

```

Epoch 2

```

-----
loss: 0.610377 [ 0/ 3009]
loss: 0.635636 [ 400/ 3009]
loss: 0.459467 [ 800/ 3009]
loss: 0.638018 [1200/ 3009]
loss: 0.553621 [1600/ 3009]
loss: 0.758643 [2000/ 3009]
loss: 0.648250 [2400/ 3009]
loss: 0.499889 [2800/ 3009]
Test Error:
Accuracy: 79.0%, Avg loss: 0.501841

```

Epoch 3

```
-----  
loss: 0.425109 [ 0/ 3009]  
loss: 0.521330 [ 400/ 3009]  
loss: 0.658485 [ 800/ 3009]  
loss: 0.244719 [ 1200/ 3009]  
loss: 0.252357 [ 1600/ 3009]  
loss: 0.307023 [ 2000/ 3009]  
loss: 0.536315 [ 2400/ 3009]  
loss: 0.851961 [ 2800/ 3009]  
Test Error:  
Accuracy: 80.1%, Avg loss: 0.463312
```

#### Epoch 4

```
-----  
loss: 0.369217 [ 0/ 3009]  
loss: 0.232987 [ 400/ 3009]  
loss: 0.335672 [ 800/ 3009]  
loss: 0.706650 [ 1200/ 3009]  
loss: 0.287788 [ 1600/ 3009]  
loss: 0.169689 [ 2000/ 3009]  
loss: 0.183748 [ 2400/ 3009]  
loss: 1.081773 [ 2800/ 3009]  
Test Error:  
Accuracy: 80.5%, Avg loss: 0.449939
```

#### Epoch 5

```
-----  
loss: 0.619481 [ 0/ 3009]  
loss: 0.762072 [ 400/ 3009]  
loss: 0.309186 [ 800/ 3009]  
loss: 0.468054 [ 1200/ 3009]  
loss: 0.152687 [ 1600/ 3009]  
loss: 1.196430 [ 2000/ 3009]  
loss: 0.565655 [ 2400/ 3009]  
loss: 0.423156 [ 2800/ 3009]  
Test Error:  
Accuracy: 81.1%, Avg loss: 0.440605
```

#### Epoch 6

```
-----  
loss: 0.408311 [ 0/ 3009]  
loss: 0.406658 [ 400/ 3009]  
loss: 0.643183 [ 800/ 3009]  
loss: 0.680348 [ 1200/ 3009]  
loss: 0.292098 [ 1600/ 3009]  
loss: 0.885864 [ 2000/ 3009]  
loss: 0.384026 [ 2400/ 3009]  
loss: 0.291240 [ 2800/ 3009]
```



Test Error:

Accuracy: 81.5%, Avg loss: 0.429059

Epoch 7

```
-----  
loss: 0.972317 [ 0/ 3009]  
loss: 0.647492 [ 400/ 3009]  
loss: 0.287916 [ 800/ 3009]  
loss: 0.136399 [ 1200/ 3009]  
loss: 0.131110 [ 1600/ 3009]  
loss: 1.046456 [ 2000/ 3009]  
loss: 0.321376 [ 2400/ 3009]  
loss: 0.262513 [ 2800/ 3009]
```

Test Error:

Accuracy: 81.4%, Avg loss: 0.428422

Epoch 8

```
-----  
loss: 0.586174 [ 0/ 3009]  
loss: 0.280675 [ 400/ 3009]  
loss: 0.535293 [ 800/ 3009]  
loss: 0.510741 [ 1200/ 3009]  
loss: 0.490539 [ 1600/ 3009]  
loss: 0.403635 [ 2000/ 3009]  
loss: 0.292357 [ 2400/ 3009]  
loss: 0.243813 [ 2800/ 3009]
```

Test Error:

Accuracy: 81.1%, Avg loss: 0.425741

Epoch 9

```
-----  
loss: 0.693683 [ 0/ 3009]  
loss: 0.942616 [ 400/ 3009]  
loss: 0.519126 [ 800/ 3009]  
loss: 0.515375 [ 1200/ 3009]  
loss: 0.276121 [ 1600/ 3009]  
loss: 0.308702 [ 2000/ 3009]  
loss: 0.174399 [ 2400/ 3009]  
loss: 0.088865 [ 2800/ 3009]
```

Test Error:

Accuracy: 82.6%, Avg loss: 0.404838

Epoch 10

```
-----  
loss: 0.490923 [ 0/ 3009]  
loss: 0.250209 [ 400/ 3009]  
loss: 0.350554 [ 800/ 3009]  
loss: 0.569447 [ 1200/ 3009]
```

loss: 0.572317 [ 1600/ 3009]  
loss: 0.596234 [ 2000/ 3009]  
loss: 0.741490 [ 2400/ 3009]  
loss: 0.671194 [ 2800/ 3009]  
Test Error:  
Accuracy: 83.4%, Avg loss: 0.393213

Epoch 11

-----  
loss: 0.251380 [ 0/ 3009]  
loss: 0.490289 [ 400/ 3009]  
loss: 1.060655 [ 800/ 3009]  
loss: 0.565315 [ 1200/ 3009]  
loss: 0.849334 [ 1600/ 3009]  
loss: 0.175850 [ 2000/ 3009]  
loss: 0.606803 [ 2400/ 3009]  
loss: 0.099015 [ 2800/ 3009]  
Test Error:  
Accuracy: 83.3%, Avg loss: 0.387654

Epoch 12

-----  
loss: 0.413961 [ 0/ 3009]  
loss: 0.353818 [ 400/ 3009]  
loss: 0.492692 [ 800/ 3009]  
loss: 0.261699 [ 1200/ 3009]  
loss: 0.435635 [ 1600/ 3009]  
loss: 0.699563 [ 2000/ 3009]  
loss: 0.286699 [ 2400/ 3009]  
loss: 0.353499 [ 2800/ 3009]  
Test Error:  
Accuracy: 82.3%, Avg loss: 0.394087

Epoch 13

-----  
loss: 0.739688 [ 0/ 3009]  
loss: 0.808778 [ 400/ 3009]  
loss: 0.371513 [ 800/ 3009]  
loss: 0.347069 [ 1200/ 3009]  
loss: 0.914741 [ 1600/ 3009]  
loss: 0.343169 [ 2000/ 3009]  
loss: 0.199363 [ 2400/ 3009]  
loss: 0.571568 [ 2800/ 3009]  
Test Error:  
Accuracy: 83.9%, Avg loss: 0.373210

Epoch 14

-----

```
loss: 0.242214 [ 0/ 3009]
loss: 0.529473 [ 400/ 3009]
loss: 0.110571 [ 800/ 3009]
loss: 0.219881 [ 1200/ 3009]
loss: 0.605801 [ 1600/ 3009]
loss: 0.729178 [ 2000/ 3009]
loss: 0.451821 [ 2400/ 3009]
loss: 0.042259 [ 2800/ 3009]
Test Error:
  Accuracy: 82.7%, Avg loss: 0.387483
```

Epoch 15

```
-----
loss: 0.652645 [ 0/ 3009]
loss: 0.660620 [ 400/ 3009]
loss: 0.151364 [ 800/ 3009]
loss: 0.585606 [ 1200/ 3009]
loss: 0.187032 [ 1600/ 3009]
loss: 0.118159 [ 2000/ 3009]
loss: 0.308171 [ 2400/ 3009]
loss: 0.441323 [ 2800/ 3009]
Test Error:
  Accuracy: 85.4%, Avg loss: 0.356562
```

Epoch 16

```
-----
loss: 0.414405 [ 0/ 3009]
loss: 0.385131 [ 400/ 3009]
loss: 0.459647 [ 800/ 3009]
loss: 0.491250 [ 1200/ 3009]
loss: 0.252520 [ 1600/ 3009]
loss: 0.558517 [ 2000/ 3009]
loss: 0.276253 [ 2400/ 3009]
loss: 0.214704 [ 2800/ 3009]
Test Error:
  Accuracy: 85.1%, Avg loss: 0.347331
```

Epoch 17

```
-----
loss: 0.311998 [ 0/ 3009]
loss: 0.367243 [ 400/ 3009]
loss: 0.238140 [ 800/ 3009]
loss: 0.252870 [ 1200/ 3009]
loss: 0.517217 [ 1600/ 3009]
loss: 0.154063 [ 2000/ 3009]
loss: 0.574752 [ 2400/ 3009]
loss: 0.205588 [ 2800/ 3009]
Test Error:
```

Accuracy: 84.7%, Avg loss: 0.359407

Epoch 18

```
-----  
loss: 0.200745 [ 0/ 3009]  
loss: 0.193018 [ 400/ 3009]  
loss: 0.436503 [ 800/ 3009]  
loss: 0.351975 [ 1200/ 3009]  
loss: 0.262803 [ 1600/ 3009]  
loss: 0.339808 [ 2000/ 3009]  
loss: 0.155746 [ 2400/ 3009]  
loss: 0.349147 [ 2800/ 3009]
```

Test Error:

Accuracy: 85.5%, Avg loss: 0.338214

Epoch 19

```
-----  
loss: 0.874758 [ 0/ 3009]  
loss: 0.187211 [ 400/ 3009]  
loss: 0.266102 [ 800/ 3009]  
loss: 0.684857 [ 1200/ 3009]  
loss: 0.126850 [ 1600/ 3009]  
loss: 0.405240 [ 2000/ 3009]  
loss: 0.468569 [ 2400/ 3009]  
loss: 0.170313 [ 2800/ 3009]
```

Test Error:

Accuracy: 78.6%, Avg loss: 0.438182

Epoch 20

```
-----  
loss: 0.288537 [ 0/ 3009]  
loss: 0.459627 [ 400/ 3009]  
loss: 0.376274 [ 800/ 3009]  
loss: 0.371028 [ 1200/ 3009]  
loss: 0.326232 [ 1600/ 3009]  
loss: 0.406395 [ 2000/ 3009]  
loss: 0.324307 [ 2400/ 3009]  
loss: 0.169120 [ 2800/ 3009]
```

Test Error:

Accuracy: 86.3%, Avg loss: 0.333102

Epoch 21

```
-----  
loss: 0.579771 [ 0/ 3009]  
loss: 0.043862 [ 400/ 3009]  
loss: 0.373217 [ 800/ 3009]  
loss: 0.137673 [ 1200/ 3009]  
loss: 0.431986 [ 1600/ 3009]
```

loss: 0.180526 [ 2000/ 3009]  
loss: 0.343980 [ 2400/ 3009]  
loss: 0.182292 [ 2800/ 3009]  
Test Error:  
Accuracy: 87.0%, Avg loss: 0.304675

Epoch 22

-----  
loss: 0.243134 [ 0/ 3009]  
loss: 0.519196 [ 400/ 3009]  
loss: 0.151192 [ 800/ 3009]  
loss: 0.814512 [ 1200/ 3009]  
loss: 0.336175 [ 1600/ 3009]  
loss: 0.184668 [ 2000/ 3009]  
loss: 0.442557 [ 2400/ 3009]  
loss: 0.559415 [ 2800/ 3009]  
Test Error:  
Accuracy: 86.9%, Avg loss: 0.315558

Epoch 23

-----  
loss: 0.241646 [ 0/ 3009]  
loss: 0.121983 [ 400/ 3009]  
loss: 0.166792 [ 800/ 3009]  
loss: 0.728975 [ 1200/ 3009]  
loss: 0.853758 [ 1600/ 3009]  
loss: 0.247208 [ 2000/ 3009]  
loss: 0.210000 [ 2400/ 3009]  
loss: 0.575594 [ 2800/ 3009]  
Test Error:  
Accuracy: 87.5%, Avg loss: 0.287921

Epoch 24

-----  
loss: 0.204562 [ 0/ 3009]  
loss: 0.636178 [ 400/ 3009]  
loss: 0.353221 [ 800/ 3009]  
loss: 0.145395 [ 1200/ 3009]  
loss: 0.019980 [ 1600/ 3009]  
loss: 0.133911 [ 2000/ 3009]  
loss: 0.529071 [ 2400/ 3009]  
loss: 0.679501 [ 2800/ 3009]  
Test Error:  
Accuracy: 87.4%, Avg loss: 0.283882

Epoch 25

-----  
loss: 0.368221 [ 0/ 3009]

```
loss: 0.365529 [ 400/ 3009]
loss: 0.126380 [ 800/ 3009]
loss: 0.257044 [ 1200/ 3009]
loss: 0.531492 [ 1600/ 3009]
loss: 0.546171 [ 2000/ 3009]
loss: 0.077782 [ 2400/ 3009]
loss: 0.156337 [ 2800/ 3009]
Test Error:
  Accuracy: 88.6%, Avg loss: 0.291231
```

Epoch 26

```
-----
loss: 0.077332 [ 0/ 3009]
loss: 0.442437 [ 400/ 3009]
loss: 0.173262 [ 800/ 3009]
loss: 0.342794 [ 1200/ 3009]
loss: 0.190451 [ 1600/ 3009]
loss: 0.252560 [ 2000/ 3009]
loss: 0.106319 [ 2400/ 3009]
loss: 0.384746 [ 2800/ 3009]
Test Error:
  Accuracy: 82.6%, Avg loss: 0.379764
```

Epoch 27

```
-----
loss: 0.043573 [ 0/ 3009]
loss: 0.257958 [ 400/ 3009]
loss: 0.113562 [ 800/ 3009]
loss: 0.116501 [ 1200/ 3009]
loss: 0.608645 [ 1600/ 3009]
loss: 0.293663 [ 2000/ 3009]
loss: 0.869821 [ 2400/ 3009]
loss: 0.050862 [ 2800/ 3009]
Test Error:
  Accuracy: 88.6%, Avg loss: 0.263382
```

Epoch 28

```
-----
loss: 0.451397 [ 0/ 3009]
loss: 0.140293 [ 400/ 3009]
loss: 0.103293 [ 800/ 3009]
loss: 0.364420 [ 1200/ 3009]
loss: 0.919895 [ 1600/ 3009]
loss: 0.387272 [ 2000/ 3009]
loss: 0.113433 [ 2400/ 3009]
loss: 0.396482 [ 2800/ 3009]
Test Error:
  Accuracy: 87.8%, Avg loss: 0.287503
```

Epoch 29

```
-----  
loss: 0.240750 [ 0/ 3009]  
loss: 0.135038 [ 400/ 3009]  
loss: 0.136483 [ 800/ 3009]  
loss: 0.211313 [ 1200/ 3009]  
loss: 0.404109 [ 1600/ 3009]  
loss: 0.102361 [ 2000/ 3009]  
loss: 0.626404 [ 2400/ 3009]  
loss: 0.308952 [ 2800/ 3009]
```

Test Error:

Accuracy: 86.9%, Avg loss: 0.303608

Epoch 30

```
-----  
loss: 0.325404 [ 0/ 3009]  
loss: 0.717290 [ 400/ 3009]  
loss: 0.066294 [ 800/ 3009]  
loss: 0.311741 [ 1200/ 3009]  
loss: 0.263317 [ 1600/ 3009]  
loss: 0.092243 [ 2000/ 3009]  
loss: 0.045577 [ 2400/ 3009]  
loss: 0.319441 [ 2800/ 3009]
```

Test Error:

Accuracy: 84.9%, Avg loss: 0.320189

Epoch 31

```
-----  
loss: 0.534368 [ 0/ 3009]  
loss: 0.183639 [ 400/ 3009]  
loss: 0.370600 [ 800/ 3009]  
loss: 0.068155 [ 1200/ 3009]  
loss: 0.654767 [ 1600/ 3009]  
loss: 0.113649 [ 2000/ 3009]  
loss: 0.090134 [ 2400/ 3009]  
loss: 0.344198 [ 2800/ 3009]
```

Test Error:

Accuracy: 90.6%, Avg loss: 0.242058

Epoch 32

```
-----  
loss: 0.574545 [ 0/ 3009]  
loss: 0.063914 [ 400/ 3009]  
loss: 0.198321 [ 800/ 3009]  
loss: 0.197522 [ 1200/ 3009]  
loss: 0.062481 [ 1600/ 3009]  
loss: 0.398479 [ 2000/ 3009]
```

loss: 0.067294 [ 2400/ 3009]  
loss: 0.121848 [ 2800/ 3009]  
Test Error:  
Accuracy: 90.3%, Avg loss: 0.229504

Epoch 33

-----  
loss: 0.155469 [ 0/ 3009]  
loss: 0.525266 [ 400/ 3009]  
loss: 0.061572 [ 800/ 3009]  
loss: 0.099578 [ 1200/ 3009]  
loss: 0.167400 [ 1600/ 3009]  
loss: 0.133902 [ 2000/ 3009]  
loss: 0.022538 [ 2400/ 3009]  
loss: 0.120725 [ 2800/ 3009]  
Test Error:  
Accuracy: 88.4%, Avg loss: 0.273928

Epoch 34

-----  
loss: 0.380336 [ 0/ 3009]  
loss: 0.092845 [ 400/ 3009]  
loss: 0.227845 [ 800/ 3009]  
loss: 0.403331 [ 1200/ 3009]  
loss: 0.419965 [ 1600/ 3009]  
loss: 0.309798 [ 2000/ 3009]  
loss: 0.153747 [ 2400/ 3009]  
loss: 0.248192 [ 2800/ 3009]  
Test Error:  
Accuracy: 81.3%, Avg loss: 0.424367

Epoch 35

-----  
loss: 0.019522 [ 0/ 3009]  
loss: 0.033974 [ 400/ 3009]  
loss: 0.030182 [ 800/ 3009]  
loss: 0.490474 [ 1200/ 3009]  
loss: 0.404290 [ 1600/ 3009]  
loss: 0.031227 [ 2000/ 3009]  
loss: 0.060337 [ 2400/ 3009]  
loss: 0.035989 [ 2800/ 3009]  
Test Error:  
Accuracy: 89.0%, Avg loss: 0.247775

Epoch 36

-----  
loss: 0.136393 [ 0/ 3009]  
loss: 0.434957 [ 400/ 3009]



```
loss: 0.068037 [ 800/ 3009]
loss: 0.250394 [ 1200/ 3009]
loss: 0.379639 [ 1600/ 3009]
loss: 0.154445 [ 2000/ 3009]
loss: 0.421805 [ 2400/ 3009]
loss: 0.107628 [ 2800/ 3009]
Test Error:
  Accuracy: 90.6%, Avg loss: 0.232820
```

Epoch 37

```
-----
loss: 0.214086 [ 0/ 3009]
loss: 0.256247 [ 400/ 3009]
loss: 0.495287 [ 800/ 3009]
loss: 0.031020 [ 1200/ 3009]
loss: 0.155265 [ 1600/ 3009]
loss: 0.371716 [ 2000/ 3009]
loss: 0.227169 [ 2400/ 3009]
loss: 0.337771 [ 2800/ 3009]
Test Error:
  Accuracy: 90.8%, Avg loss: 0.216823
```

Epoch 38

```
-----
loss: 0.096700 [ 0/ 3009]
loss: 0.144805 [ 400/ 3009]
loss: 0.208566 [ 800/ 3009]
loss: 0.069215 [ 1200/ 3009]
loss: 0.052309 [ 1600/ 3009]
loss: 0.520875 [ 2000/ 3009]
loss: 0.346237 [ 2400/ 3009]
loss: 0.175983 [ 2800/ 3009]
Test Error:
  Accuracy: 91.2%, Avg loss: 0.218782
```

Epoch 39

```
-----
loss: 0.305555 [ 0/ 3009]
loss: 0.150947 [ 400/ 3009]
loss: 0.500910 [ 800/ 3009]
loss: 0.219239 [ 1200/ 3009]
loss: 0.102949 [ 1600/ 3009]
loss: 0.114788 [ 2000/ 3009]
loss: 0.164632 [ 2400/ 3009]
loss: 0.062595 [ 2800/ 3009]
Test Error:
  Accuracy: 91.2%, Avg loss: 0.226254
```

Epoch 40

```
-----  
loss: 0.006157 [ 0/ 3009]  
loss: 0.191176 [ 400/ 3009]  
loss: 0.359908 [ 800/ 3009]  
loss: 0.117878 [ 1200/ 3009]  
loss: 0.083638 [ 1600/ 3009]  
loss: 0.127077 [ 2000/ 3009]  
loss: 0.101555 [ 2400/ 3009]  
loss: 0.026675 [ 2800/ 3009]
```

Test Error:

Accuracy: 91.4%, Avg loss: 0.209517

Epoch 41

```
-----  
loss: 0.461294 [ 0/ 3009]  
loss: 0.113928 [ 400/ 3009]  
loss: 0.074724 [ 800/ 3009]  
loss: 0.534803 [ 1200/ 3009]  
loss: 0.017591 [ 1600/ 3009]  
loss: 0.078571 [ 2000/ 3009]  
loss: 0.137161 [ 2400/ 3009]  
loss: 0.194793 [ 2800/ 3009]
```

Test Error:

Accuracy: 91.5%, Avg loss: 0.204340

Epoch 42

```
-----  
loss: 0.009259 [ 0/ 3009]  
loss: 0.102977 [ 400/ 3009]  
loss: 0.011841 [ 800/ 3009]  
loss: 0.122268 [ 1200/ 3009]  
loss: 0.522349 [ 1600/ 3009]  
loss: 0.016986 [ 2000/ 3009]  
loss: 0.127133 [ 2400/ 3009]  
loss: 0.621686 [ 2800/ 3009]
```

Test Error:

Accuracy: 91.8%, Avg loss: 0.214583

Epoch 43

```
-----  
loss: 0.019270 [ 0/ 3009]  
loss: 0.194109 [ 400/ 3009]  
loss: 0.098511 [ 800/ 3009]  
loss: 0.510885 [ 1200/ 3009]  
loss: 0.074684 [ 1600/ 3009]  
loss: 0.028622 [ 2000/ 3009]  
loss: 0.003813 [ 2400/ 3009]
```

loss: 0.217556 [ 2800/ 3009]  
Test Error:  
Accuracy: 91.1%, Avg loss: 0.213688

Epoch 44

-----  
loss: 0.249906 [ 0/ 3009]  
loss: 0.036435 [ 400/ 3009]  
loss: 0.382848 [ 800/ 3009]  
loss: 0.138155 [ 1200/ 3009]  
loss: 0.015072 [ 1600/ 3009]  
loss: 0.183296 [ 2000/ 3009]  
loss: 0.765493 [ 2400/ 3009]  
loss: 0.002745 [ 2800/ 3009]  
Test Error:  
Accuracy: 91.6%, Avg loss: 0.192210

Epoch 45

-----  
loss: 0.048382 [ 0/ 3009]  
loss: 0.007643 [ 400/ 3009]  
loss: 0.002455 [ 800/ 3009]  
loss: 0.314233 [ 1200/ 3009]  
loss: 0.250882 [ 1600/ 3009]  
loss: 0.187964 [ 2000/ 3009]  
loss: 0.310398 [ 2400/ 3009]  
loss: 0.292651 [ 2800/ 3009]  
Test Error:  
Accuracy: 92.3%, Avg loss: 0.181595

Epoch 46

-----  
loss: 0.526552 [ 0/ 3009]  
loss: 0.350501 [ 400/ 3009]  
loss: 0.018045 [ 800/ 3009]  
loss: 0.241172 [ 1200/ 3009]  
loss: 0.024312 [ 1600/ 3009]  
loss: 0.064598 [ 2000/ 3009]  
loss: 0.066437 [ 2400/ 3009]  
loss: 0.121476 [ 2800/ 3009]  
Test Error:  
Accuracy: 92.2%, Avg loss: 0.192466

Epoch 47

-----  
loss: 0.237301 [ 0/ 3009]  
loss: 0.018461 [ 400/ 3009]  
loss: 0.023916 [ 800/ 3009]

loss: 0.106190 [ 1200/ 3009]  
loss: 0.033839 [ 1600/ 3009]  
loss: 0.025127 [ 2000/ 3009]  
loss: 0.095227 [ 2400/ 3009]  
loss: 0.449622 [ 2800/ 3009]  
Test Error:  
Accuracy: 92.0%, Avg loss: 0.198553

Epoch 48

-----  
loss: 0.004209 [ 0/ 3009]  
loss: 0.232386 [ 400/ 3009]  
loss: 0.401816 [ 800/ 3009]  
loss: 0.320880 [ 1200/ 3009]  
loss: 0.032645 [ 1600/ 3009]  
loss: 0.070538 [ 2000/ 3009]  
loss: 0.078433 [ 2400/ 3009]  
loss: 0.031259 [ 2800/ 3009]  
Test Error:  
Accuracy: 93.1%, Avg loss: 0.176062

Epoch 49

-----  
loss: 0.280083 [ 0/ 3009]  
loss: 0.034082 [ 400/ 3009]  
loss: 0.059922 [ 800/ 3009]  
loss: 0.024797 [ 1200/ 3009]  
loss: 0.078281 [ 1600/ 3009]  
loss: 0.443106 [ 2000/ 3009]  
loss: 0.164287 [ 2400/ 3009]  
loss: 0.131497 [ 2800/ 3009]  
Test Error:  
Accuracy: 90.8%, Avg loss: 0.238801

Epoch 50

-----  
loss: 0.008298 [ 0/ 3009]  
loss: 0.043125 [ 400/ 3009]  
loss: 0.075283 [ 800/ 3009]  
loss: 0.077505 [ 1200/ 3009]  
loss: 0.283944 [ 1600/ 3009]  
loss: 0.121972 [ 2000/ 3009]  
loss: 0.008125 [ 2400/ 3009]  
loss: 0.080740 [ 2800/ 3009]  
Test Error:  
Accuracy: 91.5%, Avg loss: 0.191261

Epoch 51

```
-----  
loss: 0.044427 [ 0/ 3009]  
loss: 0.186360 [ 400/ 3009]  
loss: 0.002991 [ 800/ 3009]  
loss: 0.129555 [ 1200/ 3009]  
loss: 0.063185 [ 1600/ 3009]  
loss: 0.137349 [ 2000/ 3009]  
loss: 0.455101 [ 2400/ 3009]  
loss: 0.053374 [ 2800/ 3009]  
Test Error:  
Accuracy: 92.2%, Avg loss: 0.172863
```

Epoch 52

```
-----  
loss: 0.143869 [ 0/ 3009]  
loss: 0.026640 [ 400/ 3009]  
loss: 0.060242 [ 800/ 3009]  
loss: 0.087954 [ 1200/ 3009]  
loss: 0.110793 [ 1600/ 3009]  
loss: 0.133597 [ 2000/ 3009]  
loss: 0.013762 [ 2400/ 3009]  
loss: 0.129186 [ 2800/ 3009]  
Test Error:  
Accuracy: 92.6%, Avg loss: 0.195690
```

Epoch 53

```
-----  
loss: 0.627114 [ 0/ 3009]  
loss: 0.474869 [ 400/ 3009]  
loss: 0.213273 [ 800/ 3009]  
loss: 0.040031 [ 1200/ 3009]  
loss: 0.007852 [ 1600/ 3009]  
loss: 0.031571 [ 2000/ 3009]  
loss: 0.001123 [ 2400/ 3009]  
loss: 0.028975 [ 2800/ 3009]  
Test Error:  
Accuracy: 92.0%, Avg loss: 0.170234
```

Epoch 54

```
-----  
loss: 0.074014 [ 0/ 3009]  
loss: 0.036522 [ 400/ 3009]  
loss: 0.025333 [ 800/ 3009]  
loss: 0.024615 [ 1200/ 3009]  
loss: 0.063585 [ 1600/ 3009]  
loss: 0.166636 [ 2000/ 3009]  
loss: 0.134701 [ 2400/ 3009]  
loss: 0.147738 [ 2800/ 3009]
```

Test Error:

Accuracy: 92.8%, Avg loss: 0.177144

Epoch 55

```
-----  
loss: 0.035748 [ 0/ 3009]  
loss: 0.011622 [ 400/ 3009]  
loss: 0.015642 [ 800/ 3009]  
loss: 0.351276 [ 1200/ 3009]  
loss: 0.360488 [ 1600/ 3009]  
loss: 0.028831 [ 2000/ 3009]  
loss: 0.750125 [ 2400/ 3009]  
loss: 0.006516 [ 2800/ 3009]
```

Test Error:

Accuracy: 92.3%, Avg loss: 0.195945

Epoch 56

```
-----  
loss: 0.032789 [ 0/ 3009]  
loss: 0.077465 [ 400/ 3009]  
loss: 0.373449 [ 800/ 3009]  
loss: 0.178403 [ 1200/ 3009]  
loss: 0.041034 [ 1600/ 3009]  
loss: 0.272264 [ 2000/ 3009]  
loss: 0.090671 [ 2400/ 3009]  
loss: 0.478518 [ 2800/ 3009]
```

Test Error:

Accuracy: 92.4%, Avg loss: 0.166679

Epoch 57

```
-----  
loss: 0.374019 [ 0/ 3009]  
loss: 0.063673 [ 400/ 3009]  
loss: 0.116436 [ 800/ 3009]  
loss: 0.044060 [ 1200/ 3009]  
loss: 0.020648 [ 1600/ 3009]  
loss: 0.317455 [ 2000/ 3009]  
loss: 0.016334 [ 2400/ 3009]  
loss: 0.136470 [ 2800/ 3009]
```

Test Error:

Accuracy: 92.4%, Avg loss: 0.173969

Epoch 58

```
-----  
loss: 0.101519 [ 0/ 3009]  
loss: 0.071241 [ 400/ 3009]  
loss: 0.150820 [ 800/ 3009]  
loss: 0.042686 [ 1200/ 3009]
```

loss: 0.021328 [ 1600/ 3009]  
loss: 0.080425 [ 2000/ 3009]  
loss: 0.288905 [ 2400/ 3009]  
loss: 0.006239 [ 2800/ 3009]  
Test Error:  
Accuracy: 93.1%, Avg loss: 0.177651

Epoch 59

-----  
loss: 0.037491 [ 0/ 3009]  
loss: 0.003153 [ 400/ 3009]  
loss: 0.253904 [ 800/ 3009]  
loss: 0.137205 [ 1200/ 3009]  
loss: 0.329658 [ 1600/ 3009]  
loss: 0.357776 [ 2000/ 3009]  
loss: 0.881824 [ 2400/ 3009]  
loss: 0.095342 [ 2800/ 3009]  
Test Error:  
Accuracy: 92.7%, Avg loss: 0.169382

Epoch 60

-----  
loss: 0.645445 [ 0/ 3009]  
loss: 0.011491 [ 400/ 3009]  
loss: 0.107692 [ 800/ 3009]  
loss: 0.069701 [ 1200/ 3009]  
loss: 0.037588 [ 1600/ 3009]  
loss: 0.721339 [ 2000/ 3009]  
loss: 0.373313 [ 2400/ 3009]  
loss: 0.023436 [ 2800/ 3009]  
Test Error:  
Accuracy: 91.1%, Avg loss: 0.243987

Epoch 61

-----  
loss: 0.006850 [ 0/ 3009]  
loss: 0.004590 [ 400/ 3009]  
loss: 0.099188 [ 800/ 3009]  
loss: 0.014713 [ 1200/ 3009]  
loss: 0.049664 [ 1600/ 3009]  
loss: 0.039915 [ 2000/ 3009]  
loss: 0.083186 [ 2400/ 3009]  
loss: 0.061643 [ 2800/ 3009]  
Test Error:  
Accuracy: 91.1%, Avg loss: 0.219652

Epoch 62

-----

```
loss: 0.143806 [ 0/ 3009]
loss: 0.165197 [ 400/ 3009]
loss: 0.039134 [ 800/ 3009]
loss: 0.225474 [ 1200/ 3009]
loss: 0.257318 [ 1600/ 3009]
loss: 0.216942 [ 2000/ 3009]
loss: 0.169730 [ 2400/ 3009]
loss: 0.000418 [ 2800/ 3009]
Test Error:
  Accuracy: 84.1%, Avg loss: 0.371053
```

Epoch 63

```
-----
loss: 0.513983 [ 0/ 3009]
loss: 0.528835 [ 400/ 3009]
loss: 0.008454 [ 800/ 3009]
loss: 0.156376 [ 1200/ 3009]
loss: 0.099866 [ 1600/ 3009]
loss: 0.076091 [ 2000/ 3009]
loss: 0.034282 [ 2400/ 3009]
loss: 0.004475 [ 2800/ 3009]
Test Error:
  Accuracy: 92.7%, Avg loss: 0.168315
```

Epoch 64

```
-----
loss: 0.254291 [ 0/ 3009]
loss: 0.268837 [ 400/ 3009]
loss: 0.113617 [ 800/ 3009]
loss: 0.462906 [ 1200/ 3009]
loss: 0.035124 [ 1600/ 3009]
loss: 0.201200 [ 2000/ 3009]
loss: 0.038360 [ 2400/ 3009]
loss: 0.038572 [ 2800/ 3009]
Test Error:
  Accuracy: 92.6%, Avg loss: 0.182171
```

Epoch 65

```
-----
loss: 0.022693 [ 0/ 3009]
loss: 0.052663 [ 400/ 3009]
loss: 0.102111 [ 800/ 3009]
loss: 0.020169 [ 1200/ 3009]
loss: 0.640791 [ 1600/ 3009]
loss: 0.272567 [ 2000/ 3009]
loss: 0.009876 [ 2400/ 3009]
loss: 0.074759 [ 2800/ 3009]
Test Error:
```



Accuracy: 92.7%, Avg loss: 0.203175

Epoch 66

```
-----  
loss: 0.037524 [ 0/ 3009]  
loss: 0.140557 [ 400/ 3009]  
loss: 0.022173 [ 800/ 3009]  
loss: 0.160867 [ 1200/ 3009]  
loss: 0.146776 [ 1600/ 3009]  
loss: 0.038691 [ 2000/ 3009]  
loss: 0.096679 [ 2400/ 3009]  
loss: 1.046202 [ 2800/ 3009]
```

Test Error:

Accuracy: 93.4%, Avg loss: 0.157990

Epoch 67

```
-----  
loss: 0.000903 [ 0/ 3009]  
loss: 0.461546 [ 400/ 3009]  
loss: 0.373352 [ 800/ 3009]  
loss: 0.004955 [ 1200/ 3009]  
loss: 0.205911 [ 1600/ 3009]  
loss: 0.028938 [ 2000/ 3009]  
loss: 0.041878 [ 2400/ 3009]  
loss: 0.018624 [ 2800/ 3009]
```

Test Error:

Accuracy: 92.8%, Avg loss: 0.167001

Epoch 68

```
-----  
loss: 0.115579 [ 0/ 3009]  
loss: 0.213361 [ 400/ 3009]  
loss: 0.009721 [ 800/ 3009]  
loss: 0.092563 [ 1200/ 3009]  
loss: 0.050504 [ 1600/ 3009]  
loss: 0.002960 [ 2000/ 3009]  
loss: 0.270609 [ 2400/ 3009]  
loss: 0.013008 [ 2800/ 3009]
```

Test Error:

Accuracy: 61.1%, Avg loss: 2.038542

Epoch 69

```
-----  
loss: 1.104445 [ 0/ 3009]  
loss: 0.010758 [ 400/ 3009]  
loss: 0.207830 [ 800/ 3009]  
loss: 0.376282 [ 1200/ 3009]  
loss: 0.010249 [ 1600/ 3009]
```

loss: 0.011536 [ 2000/ 3009]  
loss: 0.025457 [ 2400/ 3009]  
loss: 0.000742 [ 2800/ 3009]  
Test Error:  
Accuracy: 93.8%, Avg loss: 0.157803

Epoch 70

-----  
loss: 0.004517 [ 0/ 3009]  
loss: 0.021887 [ 400/ 3009]  
loss: 0.392092 [ 800/ 3009]  
loss: 0.021333 [ 1200/ 3009]  
loss: 0.029881 [ 1600/ 3009]  
loss: 0.014090 [ 2000/ 3009]  
loss: 0.576094 [ 2400/ 3009]  
loss: 0.016256 [ 2800/ 3009]  
Test Error:  
Accuracy: 92.2%, Avg loss: 0.196472

Epoch 71

-----  
loss: 0.015786 [ 0/ 3009]  
loss: 0.024099 [ 400/ 3009]  
loss: 0.010376 [ 800/ 3009]  
loss: 0.270794 [ 1200/ 3009]  
loss: 0.006614 [ 1600/ 3009]  
loss: 0.075882 [ 2000/ 3009]  
loss: 0.201820 [ 2400/ 3009]  
loss: 0.106887 [ 2800/ 3009]  
Test Error:  
Accuracy: 93.1%, Avg loss: 0.179289

Epoch 72

-----  
loss: 0.066107 [ 0/ 3009]  
loss: 0.044215 [ 400/ 3009]  
loss: 0.039005 [ 800/ 3009]  
loss: 0.040244 [ 1200/ 3009]  
loss: 0.112673 [ 1600/ 3009]  
loss: 0.066581 [ 2000/ 3009]  
loss: 0.014780 [ 2400/ 3009]  
loss: 0.001996 [ 2800/ 3009]  
Test Error:  
Accuracy: 93.2%, Avg loss: 0.156728

Epoch 73

-----  
loss: 0.054308 [ 0/ 3009]

```
loss: 0.135788 [ 400/ 3009]
loss: 0.488879 [ 800/ 3009]
loss: 0.058845 [ 1200/ 3009]
loss: 0.001655 [ 1600/ 3009]
loss: 0.170348 [ 2000/ 3009]
loss: 0.597990 [ 2400/ 3009]
loss: 0.012858 [ 2800/ 3009]
Test Error:
  Accuracy: 93.2%, Avg loss: 0.167063
```

Epoch 74

```
-----
loss: 0.006741 [ 0/ 3009]
loss: 0.009878 [ 400/ 3009]
loss: 0.023172 [ 800/ 3009]
loss: 0.049223 [ 1200/ 3009]
loss: 0.035407 [ 1600/ 3009]
loss: 0.281755 [ 2000/ 3009]
loss: 0.039603 [ 2400/ 3009]
loss: 0.014024 [ 2800/ 3009]
Test Error:
  Accuracy: 93.2%, Avg loss: 0.173177
```

Epoch 75

```
-----
loss: 0.069044 [ 0/ 3009]
loss: 0.013170 [ 400/ 3009]
loss: 0.075636 [ 800/ 3009]
loss: 0.023451 [ 1200/ 3009]
loss: 0.006656 [ 1600/ 3009]
loss: 0.006781 [ 2000/ 3009]
loss: 0.040001 [ 2400/ 3009]
loss: 0.026015 [ 2800/ 3009]
Test Error:
  Accuracy: 93.5%, Avg loss: 0.154842
```

Epoch 76

```
-----
loss: 0.129830 [ 0/ 3009]
loss: 0.040514 [ 400/ 3009]
loss: 0.097983 [ 800/ 3009]
loss: 0.010570 [ 1200/ 3009]
loss: 0.363309 [ 1600/ 3009]
loss: 0.000646 [ 2000/ 3009]
loss: 1.124698 [ 2400/ 3009]
loss: 0.000920 [ 2800/ 3009]
Test Error:
  Accuracy: 93.6%, Avg loss: 0.158166
```

Epoch 77

```
-----  
loss: 0.048524 [ 0/ 3009]  
loss: 0.013145 [ 400/ 3009]  
loss: 0.020278 [ 800/ 3009]  
loss: 0.033555 [ 1200/ 3009]  
loss: 0.019650 [ 1600/ 3009]  
loss: 0.019511 [ 2000/ 3009]  
loss: 0.086447 [ 2400/ 3009]  
loss: 0.087477 [ 2800/ 3009]
```

Test Error:

Accuracy: 93.8%, Avg loss: 0.156554

Epoch 78

```
-----  
loss: 0.033737 [ 0/ 3009]  
loss: 0.000792 [ 400/ 3009]  
loss: 0.252864 [ 800/ 3009]  
loss: 0.099363 [ 1200/ 3009]  
loss: 0.077244 [ 1600/ 3009]  
loss: 0.008394 [ 2000/ 3009]  
loss: 0.384961 [ 2400/ 3009]  
loss: 0.036216 [ 2800/ 3009]
```

Test Error:

Accuracy: 93.9%, Avg loss: 0.159181

Epoch 79

```
-----  
loss: 0.000724 [ 0/ 3009]  
loss: 0.014339 [ 400/ 3009]  
loss: 0.066182 [ 800/ 3009]  
loss: 0.110747 [ 1200/ 3009]  
loss: 0.059909 [ 1600/ 3009]  
loss: 0.487052 [ 2000/ 3009]  
loss: 0.352145 [ 2400/ 3009]  
loss: 0.001679 [ 2800/ 3009]
```

Test Error:

Accuracy: 90.6%, Avg loss: 0.245481

Epoch 80

```
-----  
loss: 0.031259 [ 0/ 3009]  
loss: 0.092865 [ 400/ 3009]  
loss: 0.002832 [ 800/ 3009]  
loss: 0.165064 [ 1200/ 3009]  
loss: 0.002293 [ 1600/ 3009]  
loss: 0.059774 [ 2000/ 3009]
```

loss: 0.000781 [ 2400/ 3009]  
loss: 0.154801 [ 2800/ 3009]  
Test Error:  
Accuracy: 94.4%, Avg loss: 0.159270

Epoch 81

-----  
loss: 0.024156 [ 0/ 3009]  
loss: 0.031090 [ 400/ 3009]  
loss: 0.006847 [ 800/ 3009]  
loss: 0.030935 [ 1200/ 3009]  
loss: 0.217824 [ 1600/ 3009]  
loss: 0.020595 [ 2000/ 3009]  
loss: 0.025340 [ 2400/ 3009]  
loss: 0.311222 [ 2800/ 3009]  
Test Error:  
Accuracy: 93.1%, Avg loss: 0.171679

Epoch 82

-----  
loss: 0.001649 [ 0/ 3009]  
loss: 0.043509 [ 400/ 3009]  
loss: 0.570816 [ 800/ 3009]  
loss: 0.002152 [ 1200/ 3009]  
loss: 0.066446 [ 1600/ 3009]  
loss: 0.005270 [ 2000/ 3009]  
loss: 0.026196 [ 2400/ 3009]  
loss: 0.007359 [ 2800/ 3009]  
Test Error:  
Accuracy: 92.6%, Avg loss: 0.202278

Epoch 83

-----  
loss: 0.000882 [ 0/ 3009]  
loss: 0.051283 [ 400/ 3009]  
loss: 0.266509 [ 800/ 3009]  
loss: 0.018254 [ 1200/ 3009]  
loss: 0.101562 [ 1600/ 3009]  
loss: 0.041997 [ 2000/ 3009]  
loss: 0.023880 [ 2400/ 3009]  
loss: 0.227639 [ 2800/ 3009]  
Test Error:  
Accuracy: 92.8%, Avg loss: 0.183190

Epoch 84

-----  
loss: 0.007798 [ 0/ 3009]  
loss: 0.000857 [ 400/ 3009]

```
loss: 0.034641 [ 800/ 3009]
loss: 0.003413 [ 1200/ 3009]
loss: 0.009463 [ 1600/ 3009]
loss: 0.047884 [ 2000/ 3009]
loss: 0.091985 [ 2400/ 3009]
loss: 0.044464 [ 2800/ 3009]
Test Error:
  Accuracy: 93.9%, Avg loss: 0.163869
```

Epoch 85

```
-----
loss: 0.020284 [ 0/ 3009]
loss: 0.063107 [ 400/ 3009]
loss: 0.000190 [ 800/ 3009]
loss: 0.005738 [ 1200/ 3009]
loss: 0.080214 [ 1600/ 3009]
loss: 0.008644 [ 2000/ 3009]
loss: 0.005532 [ 2400/ 3009]
loss: 0.018630 [ 2800/ 3009]
Test Error:
  Accuracy: 93.6%, Avg loss: 0.159843
```

Epoch 86

```
-----
loss: 0.002318 [ 0/ 3009]
loss: 0.005261 [ 400/ 3009]
loss: 0.208498 [ 800/ 3009]
loss: 0.096166 [ 1200/ 3009]
loss: 0.006423 [ 1600/ 3009]
loss: 0.373555 [ 2000/ 3009]
loss: 0.020800 [ 2400/ 3009]
loss: 0.013847 [ 2800/ 3009]
Test Error:
  Accuracy: 93.5%, Avg loss: 0.195176
```

Epoch 87

```
-----
loss: 0.001203 [ 0/ 3009]
loss: 0.047656 [ 400/ 3009]
loss: 0.175541 [ 800/ 3009]
loss: 0.023774 [ 1200/ 3009]
loss: 0.679213 [ 1600/ 3009]
loss: 0.689421 [ 2000/ 3009]
loss: 0.056002 [ 2400/ 3009]
loss: 0.019270 [ 2800/ 3009]
Test Error:
  Accuracy: 93.8%, Avg loss: 0.153779
```

Epoch 88

-----  
loss: 0.011208 [ 0/ 3009]  
loss: 0.009703 [ 400/ 3009]  
loss: 0.010386 [ 800/ 3009]  
loss: 0.022391 [ 1200/ 3009]  
loss: 0.008434 [ 1600/ 3009]  
loss: 0.191793 [ 2000/ 3009]  
loss: 0.623906 [ 2400/ 3009]  
loss: 0.013278 [ 2800/ 3009]

Test Error:

Accuracy: 94.3%, Avg loss: 0.150847

Epoch 89

-----  
loss: 0.018373 [ 0/ 3009]  
loss: 0.054582 [ 400/ 3009]  
loss: 0.005354 [ 800/ 3009]  
loss: 0.232555 [ 1200/ 3009]  
loss: 0.071037 [ 1600/ 3009]  
loss: 0.011367 [ 2000/ 3009]  
loss: 0.036455 [ 2400/ 3009]  
loss: 0.044183 [ 2800/ 3009]

Test Error:

Accuracy: 93.9%, Avg loss: 0.156266

Epoch 90

-----  
loss: 0.004823 [ 0/ 3009]  
loss: 0.083422 [ 400/ 3009]  
loss: 0.423318 [ 800/ 3009]  
loss: 0.013316 [ 1200/ 3009]  
loss: 0.132344 [ 1600/ 3009]  
loss: 0.111381 [ 2000/ 3009]  
loss: 0.000617 [ 2400/ 3009]  
loss: 0.016089 [ 2800/ 3009]

Test Error:

Accuracy: 93.6%, Avg loss: 0.173889

Epoch 91

-----  
loss: 0.005728 [ 0/ 3009]  
loss: 0.266697 [ 400/ 3009]  
loss: 0.001900 [ 800/ 3009]  
loss: 0.119340 [ 1200/ 3009]  
loss: 0.035948 [ 1600/ 3009]  
loss: 0.080968 [ 2000/ 3009]  
loss: 0.078530 [ 2400/ 3009]

loss: 0.166100 [ 2800/ 3009]  
Test Error:  
Accuracy: 94.3%, Avg loss: 0.143272

Epoch 92

-----  
loss: 0.001165 [ 0/ 3009]  
loss: 0.024623 [ 400/ 3009]  
loss: 0.000607 [ 800/ 3009]  
loss: 0.000455 [ 1200/ 3009]  
loss: 0.005446 [ 1600/ 3009]  
loss: 0.007006 [ 2000/ 3009]  
loss: 0.007421 [ 2400/ 3009]  
loss: 0.002075 [ 2800/ 3009]  
Test Error:  
Accuracy: 94.3%, Avg loss: 0.154272

Epoch 93

-----  
loss: 0.018297 [ 0/ 3009]  
loss: 0.003098 [ 400/ 3009]  
loss: 0.392535 [ 800/ 3009]  
loss: 0.152666 [ 1200/ 3009]  
loss: 0.308703 [ 1600/ 3009]  
loss: 0.343534 [ 2000/ 3009]  
loss: 0.028807 [ 2400/ 3009]  
loss: 0.055016 [ 2800/ 3009]  
Test Error:  
Accuracy: 94.4%, Avg loss: 0.158419

Epoch 94

-----  
loss: 0.077188 [ 0/ 3009]  
loss: 0.020270 [ 400/ 3009]  
loss: 0.002206 [ 800/ 3009]  
loss: 0.070741 [ 1200/ 3009]  
loss: 0.017653 [ 1600/ 3009]  
loss: 0.213295 [ 2000/ 3009]  
loss: 0.486937 [ 2400/ 3009]  
loss: 0.017166 [ 2800/ 3009]  
Test Error:  
Accuracy: 95.0%, Avg loss: 0.151423

Epoch 95

-----  
loss: 0.023910 [ 0/ 3009]  
loss: 0.222534 [ 400/ 3009]  
loss: 0.000244 [ 800/ 3009]



loss: 0.054321 [ 1200/ 3009]  
loss: 0.030470 [ 1600/ 3009]  
loss: 0.025719 [ 2000/ 3009]  
loss: 0.052564 [ 2400/ 3009]  
loss: 0.016367 [ 2800/ 3009]  
Test Error:  
Accuracy: 78.2%, Avg loss: 0.846202

Epoch 96

-----  
loss: 0.014835 [ 0/ 3009]  
loss: 0.210507 [ 400/ 3009]  
loss: 0.037202 [ 800/ 3009]  
loss: 0.048875 [ 1200/ 3009]  
loss: 0.494718 [ 1600/ 3009]  
loss: 0.209444 [ 2000/ 3009]  
loss: 0.001513 [ 2400/ 3009]  
loss: 0.009607 [ 2800/ 3009]  
Test Error:  
Accuracy: 94.6%, Avg loss: 0.149545

Epoch 97

-----  
loss: 0.135761 [ 0/ 3009]  
loss: 0.051289 [ 400/ 3009]  
loss: 0.003405 [ 800/ 3009]  
loss: 0.030093 [ 1200/ 3009]  
loss: 0.257032 [ 1600/ 3009]  
loss: 0.006329 [ 2000/ 3009]  
loss: 0.018746 [ 2400/ 3009]  
loss: 0.005541 [ 2800/ 3009]  
Test Error:  
Accuracy: 93.9%, Avg loss: 0.152704

Epoch 98

-----  
loss: 0.143142 [ 0/ 3009]  
loss: 0.046331 [ 400/ 3009]  
loss: 0.057589 [ 800/ 3009]  
loss: 0.016518 [ 1200/ 3009]  
loss: 0.008309 [ 1600/ 3009]  
loss: 0.004192 [ 2000/ 3009]  
loss: 0.086704 [ 2400/ 3009]  
loss: 0.001388 [ 2800/ 3009]  
Test Error:  
Accuracy: 95.5%, Avg loss: 0.146856

Epoch 99

```
-----  
loss: 0.072758 [ 0/ 3009]  
loss: 0.000313 [ 400/ 3009]  
loss: 0.196128 [ 800/ 3009]  
loss: 0.387197 [ 1200/ 3009]  
loss: 0.030790 [ 1600/ 3009]  
loss: 0.009735 [ 2000/ 3009]  
loss: 0.017582 [ 2400/ 3009]  
loss: 0.388105 [ 2800/ 3009]  
Test Error:  
Accuracy: 94.6%, Avg loss: 0.155905
```


Epoch 100

```
-----  
loss: 0.009640 [ 0/ 3009]  
loss: 0.019854 [ 400/ 3009]  
loss: 0.045112 [ 800/ 3009]  
loss: 0.009851 [ 1200/ 3009]  
loss: 0.031610 [ 1600/ 3009]  
loss: 0.015031 [ 2000/ 3009]  
loss: 0.002088 [ 2400/ 3009]  
loss: 0.003318 [ 2800/ 3009]  
Test Error:  
Accuracy: 91.5%, Avg loss: 0.234863
```

Done!

```
[60]: plt.plot(np.arange(0,100),losses)
```

```
[60]: [<matplotlib.lines.Line2D at 0x7fa27f911710>]
```



```
[88]: class NeuralNetwork2(nn.Module):
    def __init__(self):
        super(NeuralNetwork2, self).__init__()
        self.stack = nn.Sequential(
            nn.Conv2d(3,32,kernel_size=3,stride=1,padding=1),
            nn.ReLU(),
            nn.Conv2d(32,64,kernel_size=3,stride=1,padding=1),
            nn.ReLU(),
            nn.MaxPool2d(2,2),
            nn.Flatten(),
            nn.Linear(36864,512),
            nn.ReLU(),
            nn.Linear(512,24),
            nn.ReLU(),
            nn.Linear(24,2)
        )

    def forward(self, x):
        logits = self.stack(x)
        return logits
```

```
[89]: model2 = NeuralNetwork2().to(cuda)
print(model2)
```

```
NeuralNetwork2(
  (stack): Sequential(
```

```

(0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(1): ReLU()
(2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(3): ReLU()
(4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(5): Flatten(start_dim=1, end_dim=-1)
(6): Linear(in_features=36864, out_features=512, bias=True)
(7): ReLU()
(8): Linear(in_features=512, out_features=24, bias=True)
(9): ReLU()
(10): Linear(in_features=24, out_features=2, bias=True)
)
)

```

```

[90]: loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model2.parameters(), lr=learning_rate)

losses=[]
for t in range(epochs):
    print(f"Epoch {t+1}\n-----")
    train_loop(train_loader, model2, loss_fn, optimizer)
    loss = test_loop(test_loader, model2, loss_fn)
    losses.append(loss)
print("Done!")

```

Epoch 1

```

-----
loss: 0.711715 [ 0/ 3009]
loss: 0.705897 [ 400/ 3009]
loss: 0.781048 [ 800/ 3009]
loss: 0.708077 [ 1200/ 3009]
loss: 0.695257 [ 1600/ 3009]
loss: 0.691893 [ 2000/ 3009]
loss: 0.693443 [ 2400/ 3009]
loss: 0.642027 [ 2800/ 3009]
Test Error:
Accuracy: 56.2%, Avg loss: 0.673766

```

Epoch 2

```

-----
loss: 0.755350 [ 0/ 3009]
loss: 0.700148 [ 400/ 3009]
loss: 0.718120 [ 800/ 3009]
loss: 0.629210 [ 1200/ 3009]
loss: 0.644932 [ 1600/ 3009]
loss: 0.655081 [ 2000/ 3009]
loss: 0.669446 [ 2400/ 3009]

```

loss: 0.587303 [ 2800/ 3009]  
Test Error:  
Accuracy: 62.3%, Avg loss: 0.640227

#### Epoch 3

-----  
loss: 0.762679 [ 0/ 3009]  
loss: 0.626614 [ 400/ 3009]  
loss: 0.748828 [ 800/ 3009]  
loss: 0.653227 [ 1200/ 3009]  
loss: 0.678577 [ 1600/ 3009]  
loss: 0.457607 [ 2000/ 3009]  
loss: 0.869075 [ 2400/ 3009]  
loss: 0.577134 [ 2800/ 3009]  
Test Error:  
Accuracy: 76.1%, Avg loss: 0.556132

#### Epoch 4

-----  
loss: 0.693505 [ 0/ 3009]  
loss: 0.502505 [ 400/ 3009]  
loss: 0.425967 [ 800/ 3009]  
loss: 0.697888 [ 1200/ 3009]  
loss: 0.392529 [ 1600/ 3009]  
loss: 0.502759 [ 2000/ 3009]  
loss: 0.370721 [ 2400/ 3009]  
loss: 0.807017 [ 2800/ 3009]  
Test Error:  
Accuracy: 77.6%, Avg loss: 0.514657

#### Epoch 5

-----  
loss: 0.331364 [ 0/ 3009]  
loss: 0.493531 [ 400/ 3009]  
loss: 0.517693 [ 800/ 3009]  
loss: 0.477822 [ 1200/ 3009]  
loss: 0.254994 [ 1600/ 3009]  
loss: 0.399225 [ 2000/ 3009]  
loss: 0.264680 [ 2400/ 3009]  
loss: 0.351955 [ 2800/ 3009]  
Test Error:  
Accuracy: 78.4%, Avg loss: 0.509381

#### Epoch 6

-----  
loss: 0.865306 [ 0/ 3009]  
loss: 0.172952 [ 400/ 3009]  
loss: 0.841517 [ 800/ 3009]

loss: 0.846098 [ 1200/ 3009]  
loss: 0.527412 [ 1600/ 3009]  
loss: 1.547148 [ 2000/ 3009]  
loss: 0.479856 [ 2400/ 3009]  
loss: 0.296828 [ 2800/ 3009]  
Test Error:  
Accuracy: 79.8%, Avg loss: 0.472175

Epoch 7

-----  
loss: 0.177903 [ 0/ 3009]  
loss: 0.282860 [ 400/ 3009]  
loss: 0.170122 [ 800/ 3009]  
loss: 0.497031 [ 1200/ 3009]  
loss: 0.350640 [ 1600/ 3009]  
loss: 0.531191 [ 2000/ 3009]  
loss: 0.632520 [ 2400/ 3009]  
loss: 0.725240 [ 2800/ 3009]  
Test Error:  
Accuracy: 81.1%, Avg loss: 0.442011

Epoch 8

-----  
loss: 0.239665 [ 0/ 3009]  
loss: 0.748299 [ 400/ 3009]  
loss: 0.266596 [ 800/ 3009]  
loss: 0.966097 [ 1200/ 3009]  
loss: 0.254812 [ 1600/ 3009]  
loss: 0.729323 [ 2000/ 3009]  
loss: 0.138737 [ 2400/ 3009]  
loss: 0.129802 [ 2800/ 3009]  
Test Error:  
Accuracy: 81.5%, Avg loss: 0.429140

Epoch 9

-----  
loss: 0.141641 [ 0/ 3009]  
loss: 0.287038 [ 400/ 3009]  
loss: 0.919959 [ 800/ 3009]  
loss: 0.130690 [ 1200/ 3009]  
loss: 0.357816 [ 1600/ 3009]  
loss: 0.746361 [ 2000/ 3009]  
loss: 0.303294 [ 2400/ 3009]  
loss: 0.324328 [ 2800/ 3009]  
Test Error:  
Accuracy: 82.5%, Avg loss: 0.425929

Epoch 10

```
-----  
loss: 0.526454 [ 0/ 3009]  
loss: 0.211183 [ 400/ 3009]  
loss: 0.317216 [ 800/ 3009]  
loss: 0.294704 [ 1200/ 3009]  
loss: 0.219923 [ 1600/ 3009]  
loss: 0.216502 [ 2000/ 3009]  
loss: 0.247918 [ 2400/ 3009]  
loss: 0.158158 [ 2800/ 3009]  
Test Error:  
Accuracy: 83.5%, Avg loss: 0.409887
```

#### Epoch 11

```
-----  
loss: 0.108303 [ 0/ 3009]  
loss: 0.234122 [ 400/ 3009]  
loss: 0.124939 [ 800/ 3009]  
loss: 0.457307 [ 1200/ 3009]  
loss: 0.872024 [ 1600/ 3009]  
loss: 0.561566 [ 2000/ 3009]  
loss: 0.084783 [ 2400/ 3009]  
loss: 0.267610 [ 2800/ 3009]  
Test Error:  
Accuracy: 83.5%, Avg loss: 0.406972
```

#### Epoch 12

```
-----  
loss: 0.988234 [ 0/ 3009]  
loss: 0.207310 [ 400/ 3009]  
loss: 0.243810 [ 800/ 3009]  
loss: 0.210831 [ 1200/ 3009]  
loss: 0.140878 [ 1600/ 3009]  
loss: 0.580248 [ 2000/ 3009]  
loss: 0.743210 [ 2400/ 3009]  
loss: 0.205702 [ 2800/ 3009]  
Test Error:  
Accuracy: 81.5%, Avg loss: 0.440751
```

#### Epoch 13

```
-----  
loss: 0.138902 [ 0/ 3009]  
loss: 0.313473 [ 400/ 3009]  
loss: 0.896971 [ 800/ 3009]  
loss: 1.614543 [ 1200/ 3009]  
loss: 1.071629 [ 1600/ 3009]  
loss: 0.077596 [ 2000/ 3009]  
loss: 0.353230 [ 2400/ 3009]  
loss: 0.278893 [ 2800/ 3009]
```

Test Error:

Accuracy: 75.2%, Avg loss: 0.538745

Epoch 14

```
-----  
loss: 0.642798 [ 0/ 3009]  
loss: 0.885413 [ 400/ 3009]  
loss: 0.242328 [ 800/ 3009]  
loss: 0.276141 [ 1200/ 3009]  
loss: 0.490522 [ 1600/ 3009]  
loss: 0.519415 [ 2000/ 3009]  
loss: 0.796545 [ 2400/ 3009]  
loss: 0.885233 [ 2800/ 3009]
```

Test Error:

Accuracy: 75.4%, Avg loss: 0.521557

Epoch 15

```
-----  
loss: 0.052964 [ 0/ 3009]  
loss: 0.373672 [ 400/ 3009]  
loss: 0.408586 [ 800/ 3009]  
loss: 0.239320 [ 1200/ 3009]  
loss: 0.473551 [ 1600/ 3009]  
loss: 0.848234 [ 2000/ 3009]  
loss: 0.147464 [ 2400/ 3009]  
loss: 0.040510 [ 2800/ 3009]
```

Test Error:

Accuracy: 85.5%, Avg loss: 0.363575

Epoch 16

```
-----  
loss: 0.210684 [ 0/ 3009]  
loss: 0.045483 [ 400/ 3009]  
loss: 0.219792 [ 800/ 3009]  
loss: 0.464515 [ 1200/ 3009]  
loss: 0.265839 [ 1600/ 3009]  
loss: 0.173137 [ 2000/ 3009]  
loss: 0.727704 [ 2400/ 3009]  
loss: 0.196887 [ 2800/ 3009]
```

Test Error:

Accuracy: 85.3%, Avg loss: 0.362565

Epoch 17

```
-----  
loss: 0.783489 [ 0/ 3009]  
loss: 0.447749 [ 400/ 3009]  
loss: 0.249983 [ 800/ 3009]  
loss: 0.068901 [ 1200/ 3009]
```



loss: 0.469351 [ 1600/ 3009]  
loss: 0.206290 [ 2000/ 3009]  
loss: 0.185837 [ 2400/ 3009]  
loss: 0.923424 [ 2800/ 3009]  
Test Error:  
Accuracy: 85.8%, Avg loss: 0.353814

Epoch 18

-----  
loss: 0.595846 [ 0/ 3009]  
loss: 0.541249 [ 400/ 3009]  
loss: 0.294264 [ 800/ 3009]  
loss: 0.206931 [ 1200/ 3009]  
loss: 0.292013 [ 1600/ 3009]  
loss: 0.524441 [ 2000/ 3009]  
loss: 0.190251 [ 2400/ 3009]  
loss: 0.056238 [ 2800/ 3009]  
Test Error:  
Accuracy: 86.1%, Avg loss: 0.347850

Epoch 19

-----  
loss: 0.118232 [ 0/ 3009]  
loss: 0.494000 [ 400/ 3009]  
loss: 0.255106 [ 800/ 3009]  
loss: 0.146838 [ 1200/ 3009]  
loss: 0.023117 [ 1600/ 3009]  
loss: 0.115445 [ 2000/ 3009]  
loss: 0.229088 [ 2400/ 3009]  
loss: 0.110753 [ 2800/ 3009]  
Test Error:  
Accuracy: 86.3%, Avg loss: 0.350550

Epoch 20

-----  
loss: 0.556452 [ 0/ 3009]  
loss: 0.235874 [ 400/ 3009]  
loss: 0.068152 [ 800/ 3009]  
loss: 0.126434 [ 1200/ 3009]  
loss: 0.347096 [ 1600/ 3009]  
loss: 0.218006 [ 2000/ 3009]  
loss: 0.096414 [ 2400/ 3009]  
loss: 0.524077 [ 2800/ 3009]  
Test Error:  
Accuracy: 73.0%, Avg loss: 0.631061

Epoch 21

-----

```
loss: 0.896657 [ 0/ 3009]
loss: 0.783473 [ 400/ 3009]
loss: 0.457230 [ 800/ 3009]
loss: 0.177887 [ 1200/ 3009]
loss: 0.470935 [ 1600/ 3009]
loss: 0.225666 [ 2000/ 3009]
loss: 0.756398 [ 2400/ 3009]
loss: 0.088261 [ 2800/ 3009]
Test Error:
  Accuracy: 86.3%, Avg loss: 0.338288
```

Epoch 22

```
-----
loss: 0.146667 [ 0/ 3009]
loss: 0.717931 [ 400/ 3009]
loss: 0.242449 [ 800/ 3009]
loss: 0.175843 [ 1200/ 3009]
loss: 0.374500 [ 1600/ 3009]
loss: 1.128326 [ 2000/ 3009]
loss: 0.212621 [ 2400/ 3009]
loss: 0.181792 [ 2800/ 3009]
Test Error:
  Accuracy: 71.0%, Avg loss: 0.680837
```

Epoch 23

```
-----
loss: 1.254139 [ 0/ 3009]
loss: 0.086783 [ 400/ 3009]
loss: 0.014026 [ 800/ 3009]
loss: 0.081508 [ 1200/ 3009]
loss: 0.136452 [ 1600/ 3009]
loss: 0.699555 [ 2000/ 3009]
loss: 0.515623 [ 2400/ 3009]
loss: 0.434245 [ 2800/ 3009]
Test Error:
  Accuracy: 82.3%, Avg loss: 0.413458
```

Epoch 24

```
-----
loss: 0.155622 [ 0/ 3009]
loss: 0.037453 [ 400/ 3009]
loss: 0.150519 [ 800/ 3009]
loss: 0.353349 [ 1200/ 3009]
loss: 0.113665 [ 1600/ 3009]
loss: 0.104747 [ 2000/ 3009]
loss: 0.161026 [ 2400/ 3009]
loss: 0.670281 [ 2800/ 3009]
Test Error:
```

Accuracy: 79.0%, Avg loss: 0.477781

Epoch 25

```
-----  
loss: 0.431234 [ 0/ 3009]  
loss: 0.711103 [ 400/ 3009]  
loss: 0.160983 [ 800/ 3009]  
loss: 0.661031 [ 1200/ 3009]  
loss: 0.132485 [ 1600/ 3009]  
loss: 0.162321 [ 2000/ 3009]  
loss: 0.353385 [ 2400/ 3009]  
loss: 0.294625 [ 2800/ 3009]
```

Test Error:

Accuracy: 85.1%, Avg loss: 0.366820

Epoch 26

```
-----  
loss: 0.213271 [ 0/ 3009]  
loss: 0.020402 [ 400/ 3009]  
loss: 0.034926 [ 800/ 3009]  
loss: 0.597366 [ 1200/ 3009]  
loss: 0.405835 [ 1600/ 3009]  
loss: 0.100659 [ 2000/ 3009]  
loss: 0.100782 [ 2400/ 3009]  
loss: 0.601844 [ 2800/ 3009]
```

Test Error:

Accuracy: 86.2%, Avg loss: 0.333440

Epoch 27

```
-----  
loss: 0.039107 [ 0/ 3009]  
loss: 0.128670 [ 400/ 3009]  
loss: 0.130664 [ 800/ 3009]  
loss: 0.120917 [ 1200/ 3009]  
loss: 0.017256 [ 1600/ 3009]  
loss: 0.277486 [ 2000/ 3009]  
loss: 0.056749 [ 2400/ 3009]  
loss: 0.505215 [ 2800/ 3009]
```

Test Error:

Accuracy: 87.0%, Avg loss: 0.327390

Epoch 28

```
-----  
loss: 0.558317 [ 0/ 3009]  
loss: 0.057266 [ 400/ 3009]  
loss: 0.576435 [ 800/ 3009]  
loss: 0.370497 [ 1200/ 3009]  
loss: 0.083310 [ 1600/ 3009]
```

loss: 0.114504 [ 2000/ 3009]  
loss: 0.097639 [ 2400/ 3009]  
loss: 0.227715 [ 2800/ 3009]  
Test Error:  
Accuracy: 87.3%, Avg loss: 0.317735

Epoch 29

-----  
loss: 0.537438 [ 0/ 3009]  
loss: 0.029624 [ 400/ 3009]  
loss: 0.631468 [ 800/ 3009]  
loss: 0.058956 [ 1200/ 3009]  
loss: 0.466496 [ 1600/ 3009]  
loss: 0.600036 [ 2000/ 3009]  
loss: 0.483408 [ 2400/ 3009]  
loss: 0.712210 [ 2800/ 3009]  
Test Error:  
Accuracy: 88.2%, Avg loss: 0.317857

Epoch 30

-----  
loss: 0.254080 [ 0/ 3009]  
loss: 0.066454 [ 400/ 3009]  
loss: 0.140995 [ 800/ 3009]  
loss: 0.821423 [ 1200/ 3009]  
loss: 0.819443 [ 1600/ 3009]  
loss: 0.596985 [ 2000/ 3009]  
loss: 0.617405 [ 2400/ 3009]  
loss: 0.308478 [ 2800/ 3009]  
Test Error:  
Accuracy: 88.3%, Avg loss: 0.298088

Epoch 31

-----  
loss: 0.032066 [ 0/ 3009]  
loss: 0.294722 [ 400/ 3009]  
loss: 0.628836 [ 800/ 3009]  
loss: 0.029795 [ 1200/ 3009]  
loss: 0.086005 [ 1600/ 3009]  
loss: 0.340056 [ 2000/ 3009]  
loss: 0.195732 [ 2400/ 3009]  
loss: 0.773366 [ 2800/ 3009]  
Test Error:  
Accuracy: 87.6%, Avg loss: 0.308490

Epoch 32

-----  
loss: 0.487176 [ 0/ 3009]

```
loss: 0.512542 [ 400/ 3009]
loss: 0.143656 [ 800/ 3009]
loss: 0.343512 [1200/ 3009]
loss: 0.361330 [1600/ 3009]
loss: 0.090653 [2000/ 3009]
loss: 0.338889 [2400/ 3009]
loss: 0.651822 [2800/ 3009]
Test Error:
  Accuracy: 85.9%, Avg loss: 0.323582
```

#### Epoch 33

```
-----
loss: 0.143074 [ 0/ 3009]
loss: 0.213931 [ 400/ 3009]
loss: 0.058581 [ 800/ 3009]
loss: 0.102967 [1200/ 3009]
loss: 0.015905 [1600/ 3009]
loss: 0.113902 [2000/ 3009]
loss: 0.289573 [2400/ 3009]
loss: 0.018185 [2800/ 3009]
Test Error:
  Accuracy: 86.7%, Avg loss: 0.322337
```

#### Epoch 34

```
-----
loss: 0.013677 [ 0/ 3009]
loss: 0.027302 [ 400/ 3009]
loss: 0.176143 [ 800/ 3009]
loss: 0.006783 [1200/ 3009]
loss: 0.025181 [1600/ 3009]
loss: 0.383826 [2000/ 3009]
loss: 0.151664 [2400/ 3009]
loss: 0.047994 [2800/ 3009]
Test Error:
  Accuracy: 88.0%, Avg loss: 0.298101
```

#### Epoch 35

```
-----
loss: 0.095547 [ 0/ 3009]
loss: 0.292424 [ 400/ 3009]
loss: 0.011002 [ 800/ 3009]
loss: 0.019398 [1200/ 3009]
loss: 0.033750 [1600/ 3009]
loss: 0.058871 [2000/ 3009]
loss: 0.097687 [2400/ 3009]
loss: 0.640427 [2800/ 3009]
Test Error:
  Accuracy: 87.5%, Avg loss: 0.291619
```

Epoch 36

```
-----  
loss: 0.021856 [ 0/ 3009]  
loss: 0.049799 [ 400/ 3009]  
loss: 0.072903 [ 800/ 3009]  
loss: 0.485284 [ 1200/ 3009]  
loss: 0.549467 [ 1600/ 3009]  
loss: 0.062470 [ 2000/ 3009]  
loss: 0.016815 [ 2400/ 3009]  
loss: 0.468386 [ 2800/ 3009]
```

Test Error:

Accuracy: 80.1%, Avg loss: 0.558996

Epoch 37

```
-----  
loss: 0.027272 [ 0/ 3009]  
loss: 0.047983 [ 400/ 3009]  
loss: 0.148794 [ 800/ 3009]  
loss: 0.106693 [ 1200/ 3009]  
loss: 0.152875 [ 1600/ 3009]  
loss: 0.533970 [ 2000/ 3009]  
loss: 0.251033 [ 2400/ 3009]  
loss: 0.088258 [ 2800/ 3009]
```

Test Error:

Accuracy: 88.3%, Avg loss: 0.295869

Epoch 38

```
-----  
loss: 0.057280 [ 0/ 3009]  
loss: 0.042504 [ 400/ 3009]  
loss: 0.221389 [ 800/ 3009]  
loss: 0.316236 [ 1200/ 3009]  
loss: 0.131336 [ 1600/ 3009]  
loss: 0.053026 [ 2000/ 3009]  
loss: 0.170289 [ 2400/ 3009]  
loss: 0.661655 [ 2800/ 3009]
```

Test Error:

Accuracy: 88.3%, Avg loss: 0.276948

Epoch 39

```
-----  
loss: 0.057090 [ 0/ 3009]  
loss: 0.145102 [ 400/ 3009]  
loss: 0.060654 [ 800/ 3009]  
loss: 0.097369 [ 1200/ 3009]  
loss: 1.360102 [ 1600/ 3009]  
loss: 0.235112 [ 2000/ 3009]
```

loss: 0.065366 [ 2400/ 3009]  
loss: 0.114563 [ 2800/ 3009]  
Test Error:  
Accuracy: 88.2%, Avg loss: 0.272770

Epoch 40

-----  
loss: 0.093003 [ 0/ 3009]  
loss: 0.204854 [ 400/ 3009]  
loss: 0.228901 [ 800/ 3009]  
loss: 0.010832 [ 1200/ 3009]  
loss: 0.109464 [ 1600/ 3009]  
loss: 0.087519 [ 2000/ 3009]  
loss: 0.059529 [ 2400/ 3009]  
loss: 0.052867 [ 2800/ 3009]  
Test Error:  
Accuracy: 86.3%, Avg loss: 0.347420

Epoch 41

-----  
loss: 0.213702 [ 0/ 3009]  
loss: 0.228853 [ 400/ 3009]  
loss: 0.297126 [ 800/ 3009]  
loss: 0.321362 [ 1200/ 3009]  
loss: 0.680654 [ 1600/ 3009]  
loss: 0.066536 [ 2000/ 3009]  
loss: 0.091023 [ 2400/ 3009]  
loss: 0.305685 [ 2800/ 3009]  
Test Error:  
Accuracy: 89.0%, Avg loss: 0.273945

Epoch 42

-----  
loss: 0.006649 [ 0/ 3009]  
loss: 0.033944 [ 400/ 3009]  
loss: 0.540856 [ 800/ 3009]  
loss: 0.012566 [ 1200/ 3009]  
loss: 0.156743 [ 1600/ 3009]  
loss: 0.699563 [ 2000/ 3009]  
loss: 0.101107 [ 2400/ 3009]  
loss: 0.190124 [ 2800/ 3009]  
Test Error:  
Accuracy: 87.8%, Avg loss: 0.309860

Epoch 43

-----  
loss: 0.083549 [ 0/ 3009]  
loss: 0.020499 [ 400/ 3009]

```
loss: 0.052740 [ 800/ 3009]
loss: 0.114831 [ 1200/ 3009]
loss: 0.269933 [ 1600/ 3009]
loss: 1.102545 [ 2000/ 3009]
loss: 0.013960 [ 2400/ 3009]
loss: 0.138183 [ 2800/ 3009]
Test Error:
  Accuracy: 89.4%, Avg loss: 0.259819
```

Epoch 44

```
-----
loss: 0.559715 [ 0/ 3009]
loss: 0.066336 [ 400/ 3009]
loss: 0.133803 [ 800/ 3009]
loss: 0.102599 [ 1200/ 3009]
loss: 0.054592 [ 1600/ 3009]
loss: 0.029684 [ 2000/ 3009]
loss: 0.105189 [ 2400/ 3009]
loss: 0.148556 [ 2800/ 3009]
Test Error:
  Accuracy: 87.4%, Avg loss: 0.307553
```

Epoch 45

```
-----
loss: 0.010435 [ 0/ 3009]
loss: 0.024055 [ 400/ 3009]
loss: 0.035594 [ 800/ 3009]
loss: 0.234845 [ 1200/ 3009]
loss: 0.632556 [ 1600/ 3009]
loss: 0.119546 [ 2000/ 3009]
loss: 0.204551 [ 2400/ 3009]
loss: 0.217281 [ 2800/ 3009]
Test Error:
  Accuracy: 86.2%, Avg loss: 0.313429
```

Epoch 46

```
-----
loss: 0.002566 [ 0/ 3009]
loss: 0.050801 [ 400/ 3009]
loss: 0.689970 [ 800/ 3009]
loss: 0.027387 [ 1200/ 3009]
loss: 0.219548 [ 1600/ 3009]
loss: 0.035808 [ 2000/ 3009]
loss: 0.093705 [ 2400/ 3009]
loss: 0.119880 [ 2800/ 3009]
Test Error:
  Accuracy: 88.2%, Avg loss: 0.284301
```



Epoch 47

-----  
loss: 0.058908 [ 0/ 3009]  
loss: 0.099092 [ 400/ 3009]  
loss: 0.172905 [ 800/ 3009]  
loss: 0.132096 [ 1200/ 3009]  
loss: 0.037793 [ 1600/ 3009]  
loss: 0.006682 [ 2000/ 3009]  
loss: 0.253893 [ 2400/ 3009]  
loss: 0.010126 [ 2800/ 3009]

Test Error:

Accuracy: 86.5%, Avg loss: 0.397251

Epoch 48

-----  
loss: 0.325469 [ 0/ 3009]  
loss: 0.817043 [ 400/ 3009]  
loss: 0.627767 [ 800/ 3009]  
loss: 0.038809 [ 1200/ 3009]  
loss: 0.014622 [ 1600/ 3009]  
loss: 0.090391 [ 2000/ 3009]  
loss: 0.078769 [ 2400/ 3009]  
loss: 0.006134 [ 2800/ 3009]

Test Error:

Accuracy: 88.7%, Avg loss: 0.285309

Epoch 49

-----  
loss: 0.035623 [ 0/ 3009]  
loss: 0.216136 [ 400/ 3009]  
loss: 0.002869 [ 800/ 3009]  
loss: 0.019659 [ 1200/ 3009]  
loss: 0.012295 [ 1600/ 3009]  
loss: 0.050987 [ 2000/ 3009]  
loss: 0.014024 [ 2400/ 3009]  
loss: 0.012880 [ 2800/ 3009]

Test Error:

Accuracy: 91.4%, Avg loss: 0.212954

Epoch 50

-----  
loss: 0.102933 [ 0/ 3009]  
loss: 0.718785 [ 400/ 3009]  
loss: 0.037147 [ 800/ 3009]  
loss: 0.059148 [ 1200/ 3009]  
loss: 0.035494 [ 1600/ 3009]  
loss: 0.224930 [ 2000/ 3009]  
loss: 0.004958 [ 2400/ 3009]

loss: 0.147441 [ 2800/ 3009]  
Test Error:  
Accuracy: 90.3%, Avg loss: 0.222343

Epoch 51

-----  
loss: 0.015586 [ 0/ 3009]  
loss: 0.062767 [ 400/ 3009]  
loss: 0.004036 [ 800/ 3009]  
loss: 0.005925 [ 1200/ 3009]  
loss: 0.021460 [ 1600/ 3009]  
loss: 0.002304 [ 2000/ 3009]  
loss: 0.100230 [ 2400/ 3009]  
loss: 0.384053 [ 2800/ 3009]  
Test Error:  
Accuracy: 90.4%, Avg loss: 0.227728

Epoch 52

-----  
loss: 0.536548 [ 0/ 3009]  
loss: 0.056401 [ 400/ 3009]  
loss: 0.028797 [ 800/ 3009]  
loss: 0.231863 [ 1200/ 3009]  
loss: 0.012937 [ 1600/ 3009]  
loss: 0.036724 [ 2000/ 3009]  
loss: 0.004671 [ 2400/ 3009]  
loss: 0.086848 [ 2800/ 3009]  
Test Error:  
Accuracy: 90.0%, Avg loss: 0.235339

Epoch 53

-----  
loss: 0.180087 [ 0/ 3009]  
loss: 0.011248 [ 400/ 3009]  
loss: 0.025772 [ 800/ 3009]  
loss: 0.056170 [ 1200/ 3009]  
loss: 0.029920 [ 1600/ 3009]  
loss: 0.023130 [ 2000/ 3009]  
loss: 0.119869 [ 2400/ 3009]  
loss: 0.010935 [ 2800/ 3009]  
Test Error:  
Accuracy: 91.0%, Avg loss: 0.215204

Epoch 54

-----  
loss: 0.052787 [ 0/ 3009]  
loss: 0.159526 [ 400/ 3009]  
loss: 0.048870 [ 800/ 3009]

loss: 0.098035 [ 1200/ 3009]  
loss: 0.057270 [ 1600/ 3009]  
loss: 0.028382 [ 2000/ 3009]  
loss: 0.000872 [ 2400/ 3009]  
loss: 0.055467 [ 2800/ 3009]  
Test Error:  
Accuracy: 91.5%, Avg loss: 0.211033

Epoch 55

-----  
loss: 0.066668 [ 0/ 3009]  
loss: 0.045427 [ 400/ 3009]  
loss: 0.030634 [ 800/ 3009]  
loss: 0.000410 [ 1200/ 3009]  
loss: 0.002358 [ 1600/ 3009]  
loss: 0.490756 [ 2000/ 3009]  
loss: 0.141261 [ 2400/ 3009]  
loss: 0.005538 [ 2800/ 3009]  
Test Error:  
Accuracy: 91.5%, Avg loss: 0.216717

Epoch 56

-----  
loss: 0.189964 [ 0/ 3009]  
loss: 0.061585 [ 400/ 3009]  
loss: 0.059436 [ 800/ 3009]  
loss: 0.098524 [ 1200/ 3009]  
loss: 0.022572 [ 1600/ 3009]  
loss: 0.325416 [ 2000/ 3009]  
loss: 0.009113 [ 2400/ 3009]  
loss: 0.483147 [ 2800/ 3009]  
Test Error:  
Accuracy: 91.2%, Avg loss: 0.221143

Epoch 57

-----  
loss: 0.123602 [ 0/ 3009]  
loss: 0.012541 [ 400/ 3009]  
loss: 0.008924 [ 800/ 3009]  
loss: 0.028001 [ 1200/ 3009]  
loss: 0.041111 [ 1600/ 3009]  
loss: 0.161316 [ 2000/ 3009]  
loss: 0.439316 [ 2400/ 3009]  
loss: 0.014576 [ 2800/ 3009]  
Test Error:  
Accuracy: 91.4%, Avg loss: 0.218700

Epoch 58

```
-----  
loss: 0.001368 [ 0/ 3009]  
loss: 0.011280 [ 400/ 3009]  
loss: 0.229602 [ 800/ 3009]  
loss: 0.164365 [ 1200/ 3009]  
loss: 0.001602 [ 1600/ 3009]  
loss: 0.295999 [ 2000/ 3009]  
loss: 0.024721 [ 2400/ 3009]  
loss: 0.239548 [ 2800/ 3009]  
Test Error:  
Accuracy: 91.4%, Avg loss: 0.208826
```

Epoch 59

```
-----  
loss: 0.125604 [ 0/ 3009]  
loss: 0.014347 [ 400/ 3009]  
loss: 0.062712 [ 800/ 3009]  
loss: 0.004263 [ 1200/ 3009]  
loss: 0.101383 [ 1600/ 3009]  
loss: 0.033327 [ 2000/ 3009]  
loss: 0.007881 [ 2400/ 3009]  
loss: 0.561272 [ 2800/ 3009]  
Test Error:  
Accuracy: 91.4%, Avg loss: 0.209096
```

Epoch 60

```
-----  
loss: 0.121307 [ 0/ 3009]  
loss: 0.002635 [ 400/ 3009]  
loss: 0.071714 [ 800/ 3009]  
loss: 0.312861 [ 1200/ 3009]  
loss: 0.004490 [ 1600/ 3009]  
loss: 0.114921 [ 2000/ 3009]  
loss: 0.033915 [ 2400/ 3009]  
loss: 0.054698 [ 2800/ 3009]  
Test Error:  
Accuracy: 92.7%, Avg loss: 0.188133
```

Epoch 61

```
-----  
loss: 0.006072 [ 0/ 3009]  
loss: 0.001459 [ 400/ 3009]  
loss: 0.016800 [ 800/ 3009]  
loss: 0.021715 [ 1200/ 3009]  
loss: 0.261786 [ 1600/ 3009]  
loss: 0.008100 [ 2000/ 3009]  
loss: 0.022308 [ 2400/ 3009]  
loss: 0.006868 [ 2800/ 3009]
```

Test Error:

Accuracy: 64.8%, Avg loss: 1.186433

Epoch 62

```
-----  
loss: 1.119992 [ 0/ 3009]  
loss: 0.019007 [ 400/ 3009]  
loss: 0.124869 [ 800/ 3009]  
loss: 0.022777 [ 1200/ 3009]  
loss: 0.021165 [ 1600/ 3009]  
loss: 0.034565 [ 2000/ 3009]  
loss: 0.015211 [ 2400/ 3009]  
loss: 0.018884 [ 2800/ 3009]
```

Test Error:

Accuracy: 91.5%, Avg loss: 0.208791

Epoch 63

```
-----  
loss: 0.004318 [ 0/ 3009]  
loss: 0.029748 [ 400/ 3009]  
loss: 0.014781 [ 800/ 3009]  
loss: 0.686249 [ 1200/ 3009]  
loss: 0.323993 [ 1600/ 3009]  
loss: 0.003493 [ 2000/ 3009]  
loss: 0.132115 [ 2400/ 3009]  
loss: 0.011684 [ 2800/ 3009]
```

Test Error:

Accuracy: 93.1%, Avg loss: 0.179924

Epoch 64

```
-----  
loss: 0.028444 [ 0/ 3009]  
loss: 0.029905 [ 400/ 3009]  
loss: 0.115221 [ 800/ 3009]  
loss: 0.006342 [ 1200/ 3009]  
loss: 0.021636 [ 1600/ 3009]  
loss: 0.139053 [ 2000/ 3009]  
loss: 0.033270 [ 2400/ 3009]  
loss: 0.000578 [ 2800/ 3009]
```

Test Error:

Accuracy: 90.4%, Avg loss: 0.239954

Epoch 65

```
-----  
loss: 0.070939 [ 0/ 3009]  
loss: 0.003253 [ 400/ 3009]  
loss: 1.023592 [ 800/ 3009]  
loss: 0.010428 [ 1200/ 3009]
```

```
loss: 0.104506 [ 1600/ 3009]
loss: 0.056250 [ 2000/ 3009]
loss: 0.260863 [ 2400/ 3009]
loss: 0.002603 [ 2800/ 3009]
Test Error:
  Accuracy: 91.5%, Avg loss: 0.200238
```

Epoch 66

```
-----
loss: 0.031721 [    0/ 3009]
loss: 0.011426 [   400/ 3009]
loss: 0.025669 [   800/ 3009]
loss: 0.071854 [  1200/ 3009]
loss: 0.004346 [  1600/ 3009]
loss: 0.015355 [  2000/ 3009]
loss: 0.073710 [  2400/ 3009]
loss: 0.006099 [  2800/ 3009]
Test Error:
  Accuracy: 92.6%, Avg loss: 0.194783
```

Epoch 67

```
-----
loss: 0.056827 [    0/ 3009]
loss: 0.006316 [   400/ 3009]
loss: 0.027297 [   800/ 3009]
loss: 0.073569 [  1200/ 3009]
loss: 0.191006 [  1600/ 3009]
loss: 0.007177 [  2000/ 3009]
loss: 0.012721 [  2400/ 3009]
loss: 0.122002 [  2800/ 3009]
Test Error:
  Accuracy: 92.7%, Avg loss: 0.194939
```

Epoch 68

```
-----
loss: 0.047427 [    0/ 3009]
loss: 0.033717 [   400/ 3009]
loss: 0.046397 [   800/ 3009]
loss: 0.178196 [  1200/ 3009]
loss: 0.192990 [  1600/ 3009]
loss: 0.003153 [  2000/ 3009]
loss: 0.246383 [  2400/ 3009]
loss: 0.014472 [  2800/ 3009]
Test Error:
  Accuracy: 92.4%, Avg loss: 0.197153
```

Epoch 69

```
-----
```

```
loss: 0.012189 [ 0/ 3009]
loss: 0.008321 [ 400/ 3009]
loss: 0.015485 [ 800/ 3009]
loss: 0.155864 [ 1200/ 3009]
loss: 0.076974 [ 1600/ 3009]
loss: 0.062411 [ 2000/ 3009]
loss: 0.036185 [ 2400/ 3009]
loss: 0.098237 [ 2800/ 3009]
Test Error:
  Accuracy: 91.4%, Avg loss: 0.219328
```

Epoch 70

```
-----
loss: 0.033063 [ 0/ 3009]
loss: 0.010482 [ 400/ 3009]
loss: 1.101456 [ 800/ 3009]
loss: 0.207227 [ 1200/ 3009]
loss: 0.016984 [ 1600/ 3009]
loss: 0.060679 [ 2000/ 3009]
loss: 0.085786 [ 2400/ 3009]
loss: 0.085863 [ 2800/ 3009]
Test Error:
  Accuracy: 92.8%, Avg loss: 0.186152
```

Epoch 71

```
-----
loss: 0.014237 [ 0/ 3009]
loss: 0.008362 [ 400/ 3009]
loss: 0.059688 [ 800/ 3009]
loss: 0.019283 [ 1200/ 3009]
loss: 0.013341 [ 1600/ 3009]
loss: 0.002386 [ 2000/ 3009]
loss: 0.022993 [ 2400/ 3009]
loss: 0.063859 [ 2800/ 3009]
Test Error:
  Accuracy: 92.7%, Avg loss: 0.190775
```

Epoch 72

```
-----
loss: 0.006067 [ 0/ 3009]
loss: 0.014938 [ 400/ 3009]
loss: 0.029752 [ 800/ 3009]
loss: 0.232497 [ 1200/ 3009]
loss: 0.005524 [ 1600/ 3009]
loss: 0.039896 [ 2000/ 3009]
loss: 0.037427 [ 2400/ 3009]
loss: 0.028563 [ 2800/ 3009]
Test Error:
```

Accuracy: 92.3%, Avg loss: 0.223823

Epoch 73

```
-----  
loss: 0.023205 [ 0/ 3009]  
loss: 0.005795 [ 400/ 3009]  
loss: 0.006179 [ 800/ 3009]  
loss: 0.026254 [ 1200/ 3009]  
loss: 0.003408 [ 1600/ 3009]  
loss: 0.017294 [ 2000/ 3009]  
loss: 0.051943 [ 2400/ 3009]  
loss: 0.006927 [ 2800/ 3009]
```

Test Error:

Accuracy: 92.2%, Avg loss: 0.217712

Epoch 74

```
-----  
loss: 0.002016 [ 0/ 3009]  
loss: 0.003938 [ 400/ 3009]  
loss: 0.112243 [ 800/ 3009]  
loss: 0.045087 [ 1200/ 3009]  
loss: 0.030736 [ 1600/ 3009]  
loss: 0.027375 [ 2000/ 3009]  
loss: 0.008110 [ 2400/ 3009]  
loss: 0.085983 [ 2800/ 3009]
```

Test Error:

Accuracy: 92.0%, Avg loss: 0.210993

Epoch 75

```
-----  
loss: 0.002764 [ 0/ 3009]  
loss: 0.566190 [ 400/ 3009]  
loss: 0.015490 [ 800/ 3009]  
loss: 0.114851 [ 1200/ 3009]  
loss: 0.076364 [ 1600/ 3009]  
loss: 0.003924 [ 2000/ 3009]  
loss: 0.070145 [ 2400/ 3009]  
loss: 0.003886 [ 2800/ 3009]
```

Test Error:

Accuracy: 92.8%, Avg loss: 0.198976

Epoch 76

```
-----  
loss: 0.045867 [ 0/ 3009]  
loss: 0.028741 [ 400/ 3009]  
loss: 0.114136 [ 800/ 3009]  
loss: 0.004066 [ 1200/ 3009]  
loss: 0.004808 [ 1600/ 3009]
```



loss: 0.007023 [ 2000/ 3009]  
loss: 0.001620 [ 2400/ 3009]  
loss: 0.004190 [ 2800/ 3009]  
Test Error:  
Accuracy: 93.0%, Avg loss: 0.217826

Epoch 77

-----  
loss: 0.002527 [ 0/ 3009]  
loss: 0.008536 [ 400/ 3009]  
loss: 0.002528 [ 800/ 3009]  
loss: 0.001528 [ 1200/ 3009]  
loss: 0.036334 [ 1600/ 3009]  
loss: 0.162887 [ 2000/ 3009]  
loss: 0.018942 [ 2400/ 3009]  
loss: 0.002940 [ 2800/ 3009]  
Test Error:  
Accuracy: 92.7%, Avg loss: 0.195040

Epoch 78

-----  
loss: 0.006762 [ 0/ 3009]  
loss: 0.007818 [ 400/ 3009]  
loss: 0.084766 [ 800/ 3009]  
loss: 0.073127 [ 1200/ 3009]  
loss: 0.143585 [ 1600/ 3009]  
loss: 0.003409 [ 2000/ 3009]  
loss: 0.223054 [ 2400/ 3009]  
loss: 0.009425 [ 2800/ 3009]  
Test Error:  
Accuracy: 93.1%, Avg loss: 0.221161

Epoch 79

-----  
loss: 0.339593 [ 0/ 3009]  
loss: 0.060837 [ 400/ 3009]  
loss: 0.000541 [ 800/ 3009]  
loss: 0.015170 [ 1200/ 3009]  
loss: 0.027162 [ 1600/ 3009]  
loss: 0.007985 [ 2000/ 3009]  
loss: 0.030344 [ 2400/ 3009]  
loss: 0.011487 [ 2800/ 3009]  
Test Error:  
Accuracy: 93.0%, Avg loss: 0.185543

Epoch 80

-----  
loss: 0.229964 [ 0/ 3009]

```
loss: 0.100322 [ 400/ 3009]
loss: 0.004499 [ 800/ 3009]
loss: 0.015536 [ 1200/ 3009]
loss: 0.168483 [ 1600/ 3009]
loss: 0.040699 [ 2000/ 3009]
loss: 0.011484 [ 2400/ 3009]
loss: 0.019748 [ 2800/ 3009]
Test Error:
  Accuracy: 93.4%, Avg loss: 0.222831
```

Epoch 81

```
-----
loss: 0.053352 [ 0/ 3009]
loss: 0.010847 [ 400/ 3009]
loss: 0.075460 [ 800/ 3009]
loss: 0.017181 [ 1200/ 3009]
loss: 0.002104 [ 1600/ 3009]
loss: 0.002123 [ 2000/ 3009]
loss: 0.000692 [ 2400/ 3009]
loss: 0.013210 [ 2800/ 3009]
Test Error:
  Accuracy: 92.6%, Avg loss: 0.204206
```

Epoch 82

```
-----
loss: 0.065739 [ 0/ 3009]
loss: 0.301714 [ 400/ 3009]
loss: 0.019517 [ 800/ 3009]
loss: 0.088698 [ 1200/ 3009]
loss: 0.069003 [ 1600/ 3009]
loss: 0.002008 [ 2000/ 3009]
loss: 0.077431 [ 2400/ 3009]
loss: 0.102845 [ 2800/ 3009]
Test Error:
  Accuracy: 93.1%, Avg loss: 0.228429
```

Epoch 83

```
-----
loss: 0.952123 [ 0/ 3009]
loss: 0.039410 [ 400/ 3009]
loss: 0.009688 [ 800/ 3009]
loss: 0.033955 [ 1200/ 3009]
loss: 0.009056 [ 1600/ 3009]
loss: 0.134471 [ 2000/ 3009]
loss: 0.053211 [ 2400/ 3009]
loss: 0.012799 [ 2800/ 3009]
Test Error:
  Accuracy: 91.2%, Avg loss: 0.229423
```

Epoch 84

```
-----  
loss: 0.040025 [ 0/ 3009]  
loss: 0.001517 [ 400/ 3009]  
loss: 0.317542 [ 800/ 3009]  
loss: 0.002949 [ 1200/ 3009]  
loss: 0.011448 [ 1600/ 3009]  
loss: 0.005308 [ 2000/ 3009]  
loss: 0.023394 [ 2400/ 3009]  
loss: 0.028203 [ 2800/ 3009]
```

Test Error:

Accuracy: 93.1%, Avg loss: 0.184953

Epoch 85

```
-----  
loss: 0.204139 [ 0/ 3009]  
loss: 0.003344 [ 400/ 3009]  
loss: 0.000827 [ 800/ 3009]  
loss: 0.009044 [ 1200/ 3009]  
loss: 0.000445 [ 1600/ 3009]  
loss: 0.001391 [ 2000/ 3009]  
loss: 0.026209 [ 2400/ 3009]  
loss: 0.009114 [ 2800/ 3009]
```

Test Error:

Accuracy: 94.0%, Avg loss: 0.177088

Epoch 86

```
-----  
loss: 0.032580 [ 0/ 3009]  
loss: 0.014947 [ 400/ 3009]  
loss: 0.126224 [ 800/ 3009]  
loss: 0.057327 [ 1200/ 3009]  
loss: 0.030780 [ 1600/ 3009]  
loss: 0.001245 [ 2000/ 3009]  
loss: 0.089182 [ 2400/ 3009]  
loss: 0.004127 [ 2800/ 3009]
```

Test Error:

Accuracy: 93.9%, Avg loss: 0.195218

Epoch 87

```
-----  
loss: 0.001389 [ 0/ 3009]  
loss: 0.000323 [ 400/ 3009]  
loss: 0.022877 [ 800/ 3009]  
loss: 0.172521 [ 1200/ 3009]  
loss: 0.209852 [ 1600/ 3009]  
loss: 0.036781 [ 2000/ 3009]
```

loss: 0.002751 [ 2400/ 3009]  
loss: 0.000419 [ 2800/ 3009]  
Test Error:  
Accuracy: 89.2%, Avg loss: 0.258075

Epoch 88

-----  
loss: 0.013297 [ 0/ 3009]  
loss: 0.055379 [ 400/ 3009]  
loss: 0.087756 [ 800/ 3009]  
loss: 0.013887 [ 1200/ 3009]  
loss: 0.008034 [ 1600/ 3009]  
loss: 0.154472 [ 2000/ 3009]  
loss: 0.000505 [ 2400/ 3009]  
loss: 0.018036 [ 2800/ 3009]  
Test Error:  
Accuracy: 94.2%, Avg loss: 0.192114

Epoch 89

-----  
loss: 0.026653 [ 0/ 3009]  
loss: 0.021795 [ 400/ 3009]  
loss: 0.001559 [ 800/ 3009]  
loss: 0.195643 [ 1200/ 3009]  
loss: 0.002604 [ 1600/ 3009]  
loss: 0.047811 [ 2000/ 3009]  
loss: 0.021945 [ 2400/ 3009]  
loss: 0.074909 [ 2800/ 3009]  
Test Error:  
Accuracy: 93.4%, Avg loss: 0.190803

Epoch 90

-----  
loss: 0.011322 [ 0/ 3009]  
loss: 0.030642 [ 400/ 3009]  
loss: 0.002860 [ 800/ 3009]  
loss: 0.487855 [ 1200/ 3009]  
loss: 0.002145 [ 1600/ 3009]  
loss: 0.001327 [ 2000/ 3009]  
loss: 0.019831 [ 2400/ 3009]  
loss: 0.020152 [ 2800/ 3009]  
Test Error:  
Accuracy: 94.3%, Avg loss: 0.191063

Epoch 91

-----  
loss: 0.002234 [ 0/ 3009]  
loss: 0.002791 [ 400/ 3009]

```
loss: 0.131977 [ 800/ 3009]
loss: 0.017236 [ 1200/ 3009]
loss: 0.000345 [ 1600/ 3009]
loss: 0.154039 [ 2000/ 3009]
loss: 0.007235 [ 2400/ 3009]
loss: 0.183399 [ 2800/ 3009]
Test Error:
  Accuracy: 93.5%, Avg loss: 0.171287
```

Epoch 92

```
-----
loss: 0.017451 [ 0/ 3009]
loss: 0.007637 [ 400/ 3009]
loss: 0.019134 [ 800/ 3009]
loss: 0.000013 [ 1200/ 3009]
loss: 0.000135 [ 1600/ 3009]
loss: 0.010495 [ 2000/ 3009]
loss: 0.000343 [ 2400/ 3009]
loss: 0.010131 [ 2800/ 3009]
Test Error:
  Accuracy: 93.2%, Avg loss: 0.236074
```

Epoch 93

```
-----
loss: 0.006581 [ 0/ 3009]
loss: 0.010202 [ 400/ 3009]
loss: 0.009154 [ 800/ 3009]
loss: 0.003269 [ 1200/ 3009]
loss: 0.005554 [ 1600/ 3009]
loss: 0.000545 [ 2000/ 3009]
loss: 0.021496 [ 2400/ 3009]
loss: 0.004947 [ 2800/ 3009]
Test Error:
  Accuracy: 94.0%, Avg loss: 0.202381
```

Epoch 94

```
-----
loss: 0.009128 [ 0/ 3009]
loss: 0.022003 [ 400/ 3009]
loss: 0.010909 [ 800/ 3009]
loss: 0.033784 [ 1200/ 3009]
loss: 0.009809 [ 1600/ 3009]
loss: 0.007817 [ 2000/ 3009]
loss: 0.086461 [ 2400/ 3009]
loss: 0.256660 [ 2800/ 3009]
Test Error:
  Accuracy: 94.0%, Avg loss: 0.196094
```

Epoch 95

```
-----  
loss: 0.000411 [ 0/ 3009]  
loss: 0.133758 [ 400/ 3009]  
loss: 0.041080 [ 800/ 3009]  
loss: 0.000756 [ 1200/ 3009]  
loss: 0.000638 [ 1600/ 3009]  
loss: 0.972800 [ 2000/ 3009]  
loss: 0.016895 [ 2400/ 3009]  
loss: 0.000064 [ 2800/ 3009]
```

Test Error:

Accuracy: 93.5%, Avg loss: 0.231315

Epoch 96

```
-----  
loss: 0.000266 [ 0/ 3009]  
loss: 0.043275 [ 400/ 3009]  
loss: 0.000550 [ 800/ 3009]  
loss: 0.087086 [ 1200/ 3009]  
loss: 0.002070 [ 1600/ 3009]  
loss: 0.005622 [ 2000/ 3009]  
loss: 0.008201 [ 2400/ 3009]  
loss: 0.001006 [ 2800/ 3009]
```

Test Error:

Accuracy: 93.9%, Avg loss: 0.185028

Epoch 97

```
-----  
loss: 0.000024 [ 0/ 3009]  
loss: 0.034057 [ 400/ 3009]  
loss: 0.019065 [ 800/ 3009]  
loss: 0.000641 [ 1200/ 3009]  
loss: 0.180563 [ 1600/ 3009]  
loss: 0.002511 [ 2000/ 3009]  
loss: 0.000235 [ 2400/ 3009]  
loss: 0.001564 [ 2800/ 3009]
```

Test Error:

Accuracy: 93.8%, Avg loss: 0.182932

Epoch 98

```
-----  
loss: 0.181693 [ 0/ 3009]  
loss: 0.000027 [ 400/ 3009]  
loss: 0.000437 [ 800/ 3009]  
loss: 0.003232 [ 1200/ 3009]  
loss: 0.330453 [ 1600/ 3009]  
loss: 0.000859 [ 2000/ 3009]  
loss: 0.066065 [ 2400/ 3009]
```

```
loss: 0.058993 [ 2800/ 3009]
Test Error:
  Accuracy: 94.2%, Avg loss: 0.203471
```

Epoch 99

```
-----
loss: 0.001065 [    0/ 3009]
loss: 0.225944 [  400/ 3009]
loss: 0.019959 [  800/ 3009]
loss: 0.003069 [ 1200/ 3009]
loss: 0.067013 [ 1600/ 3009]
loss: 0.005908 [ 2000/ 3009]
loss: 0.000709 [ 2400/ 3009]
loss: 0.002509 [ 2800/ 3009]
Test Error:
  Accuracy: 94.2%, Avg loss: 0.222837
```

Epoch 100


```
-----
loss: 0.000622 [    0/ 3009]
loss: 0.011569 [  400/ 3009]
loss: 0.002145 [  800/ 3009]
loss: 0.001567 [ 1200/ 3009]
loss: 0.008113 [ 1600/ 3009]
loss: 0.300387 [ 2000/ 3009]
loss: 0.000228 [ 2400/ 3009]
loss: 0.000766 [ 2800/ 3009]
Test Error:
  Accuracy: 93.4%, Avg loss: 0.198401
```

Done!

```
[81]: 
```

```
[91]: plt.plot(np.arange(0,100),losses)
```

```
[91]: [<matplotlib.lines.Line2D at 0x7fa297297b50>]
```



```
[92]: # accuracy is better. and in first one the model is overfitted and it will not
      ↪ get better accuracy
      # but in second one as time goes, it is possible that it gets better accuracy
      ↪ and lower loss.
```

```
[ ]:
```